

Continuous-Time Costly Reversibility in Mean Field: A KS-Free Master-Equation Formulation, Derivations, and Computation

Self-contained derivation and implementation notes

September 8, 2025

Abstract

This paper derives and explains a continuous-time, mean-field (master-equation) formulation of Zhang’s costly-reversibility model. The approach is *Krusell–Smith (KS)-free*: aggregation enters through the inverse-demand dependence $P(Y(m, x))$ within the Hamiltonian, while strategic interaction across firms is encoded via the Lions derivative in the master equation. We fix primitives and state minimal boundary and regularity conditions; we then present two computational routes: (i) a stationary HJB–FP fixed point, and (ii) direct collocation of the stationary master PDE. Both routes are implementable with standard, monotone PDE schemes or modern function approximation (e.g., kernel/DeepSets representations for measures).

A central message is that the mean-field structure clarifies aggregation: the only economy-wide wedge in the firm problem is the product of the firm’s own output and the slope of inverse demand evaluated at aggregate output. Under isoelastic demand, this wedge reduces to a scalar multiple of the firm’s output. This provides a clean decomposition between *private marginal value of capital* (through the Hamiltonian) and *general-equilibrium feedback* (through the price externality). We work *conditional on the aggregate state x* , which removes common-noise second-order measure terms in the stationary master equation; Appendix C briefly outlines how those terms arise in the full common-noise setting.

We provide compact verification diagnostics (Euler and distributional residuals), explicit boundary conditions at $k = 0$ (reflecting), and growth/integrability conditions that guarantee all terms are finite. A small pseudo-JAX template illustrates how to evaluate the master-equation residual with an empirical measure. Throughout, we connect the construction to the canonical MFG literature for existence, uniqueness, and equivalence of the HJB–FP and master formulations.

Contents

Executive Summary / Cheat-Sheet (One Page)

Pedagogical Insight: Economic Intuition & Context

Primitives. Firms hold capital $k \geq 0$ and idiosyncratic productivity z . The aggregate state x shifts demand and marginal revenue. Technology is $q = e^{x+z}k^\alpha$ with $\alpha \in (0, 1)$. Inverse demand is $P(Y)$ with slope $P'(Y) < 0$, where $Y = \int e^{x+z}k^\alpha m(\cdot, dk, \cdot, dz)$. Capital follows $dk = (i - \delta k), dt$ with asymmetric, convex costs $h(i, k)$. Dividends are $\pi = P(Y)e^{x+z}k^\alpha - i - h(i, k) - f$. Shocks evolve in z and x with generators L_z, L_x . Discounting uses $r(x)$ (or constant ρ).

Core equations. Value $V(k, z, x; m)$, master value $U(k, z, x, m)$.

- **Stationary HJB:** $r(x)V = \max_i \{\pi + V_k(i - \delta k) + L_- zV + L_- xV\}$.
- **Kolmogorov–Forward (FP):** $\partial_t m = -\partial_k [(i^* - \delta k)m] + L_- z^* m$. Stationary: $\partial_t m = 0$.
- **Stationary Master Equation:** own-firm HJB terms + population-transport integrals of $D_- mU$.

Isoelastic simplification. For $P(Y) = Y^{-\eta}$, we have

$$Y P'(Y) = -\eta P(Y),$$

and therefore

$$\int \delta_m \pi, dm = -\eta P(Y) e^{x+z} k^\alpha.$$

Two solution routes.

A. HJB–FP fixed point (robust):

- 0.1. Fix x (grid/invariant law). Guess m .
- 0.2. Compute $Y, P(Y)$. Solve HJB $\Rightarrow i^*$.
- 0.3. Solve stationary FP for m' . Update $m \leftarrow m'$.

B. Direct master-PDE collocation (KS-free):

- 0.1. Parameterize U and $D_- mU$ (DeepSets/kernel for measures).
- 0.2. Build (ME) residual on empirical m (no separate externality term; price dependence enters via the Hamiltonian).
- 0.3. Penalize KKT/boundaries; recover i^* from the Hamiltonian; validate by Route A.

Diagnostics. Euler residuals for HJB, mass-balance for FP, and full ME residual. Use monotone stencils in k (upwinding) and conservative fluxes at $k = 0$.




Figure 1: Investment policy $i^*(p)$ under asymmetric adjustment costs (schematic with $k = 1$, $\phi_+ = 1$, $\phi_- = 3$).

Pedagogical Insight: Economic Intuition & Context

Recap — HJB.

- Policy is piecewise linear in $p = V_k$ with a kink at 1.
- Hamiltonian is convex in p ; envelope gives $\partial_p \mathcal{H} = i^*$.
- Reflecting boundary enforces $i^*(0, \cdot) \geq 0$ and $U_k(0, \cdot) \leq 1$.

1 Notation and Acronyms

Acronyms used in text: HJB, FP, ME, MFG, SDF, KKT, KS, RCE, TFP, CES, W2, FVM, SL.

2 Primitives and Assumptions

Assumption 2.1: Model specification; used verbatim

- (i) **Firm states:** $k \in \mathbb{R}_+$, $z \in \mathbb{R}$. **Aggregate state:** $x \in \mathbb{R}$. **Population law:** $m \in \mathcal{P}(\mathbb{R}_+ \times \mathbb{R})$.
 - (ii) **Technology:** $q(k, z, x) = e^{x+z} k^\alpha$, $\alpha \in (0, 1)$.
 - (iii) **Product market:** $P = P(Y)$ with $Y(m, x) = \int e^{x+z} k^\alpha m(\cdot, dz)$, $P'(\cdot) < 0$.
 - (iv) **Capital law:** $dk_t = (i_t - \delta k_t) dt$, $i \in \mathbb{R}$.
 - (v) **Irreversibility/adjustment:** h convex and asymmetric,
- $$h(i, k) = \begin{cases} \frac{\phi_+}{2} \frac{i^2}{k}, & i \geq 0, \\ \frac{\phi_-}{2} \frac{i^2}{k}, & i < 0, \phi_- > \phi_+. \end{cases}$$
- (vi) **Dividends:** $\pi(k, i, z, x, m) = P(Y(m, x)) e^{x+z} k^\alpha - i - h(i, k) - f$.
 - (vii) **Shocks:** $dz_t = \mu_z(z_t) dt + \sigma_z dB_t$, $dx_t = \mu_x(x_t) dt + \sigma_x dB_t$ (independent).
 - (viii) **Discounting:** short rate $r(x)$ (or constant ρ).
 - (ix) **Generators:** for smooth u ,

$$L_- z u = \mu_z(z) u_z + \frac{1}{2} \sigma_z^2 u_{zz}, \quad L_- x u = \mu_x(x) u_x + \frac{1}{2} \sigma_x^2 u_{xx}.$$

Assumption 2.2: Minimal regularity/boundary

- (a) $h(\cdot, k)$ convex, lower semicontinuous; $k \mapsto h(i, k)$ measurable with $h(i, k) \geq 0$ and $h(i, k) \geq c i^2 / k$ for some $c > 0$ on $k > 0$. The asymmetry $\phi_- > \phi_+$ holds.
- (b) P Lipschitz on compact sets with $P' < 0$; $P(Y)$ and $Y(m, x)$ finite for admissible m .

- (c) μ_z, μ_x locally Lipschitz; $\sigma_z, \sigma_x \geq 0$ constants.
- (d) *Boundary at $k = 0$* : reflecting; feasible controls satisfy $i^*(0, \cdot) \geq 0$; and $U_k(0, \cdot) \leq 1$.
- (e) *Growth*: $U(k, z, x, m) = O(k)$ as $k \rightarrow \infty$.
- (f) *Integrability*: m integrates k^α and $1/k$ wherever they appear.

Pedagogical Insight: Economic Intuition & Context

Economic reading. The convex asymmetry $\phi_- > \phi_+$ produces *investment bands*: small changes in the shadow value V_k around the frictionless cutoff 1 generate very different investment responses on the two sides of the kink. Aggregation operates through Y only, and the inverse-demand slope $P'(Y)$ is the sole channel through which the cross-section affects an individual firm's HJB. The reflecting boundary at $k = 0$ formalizes limited liability and the irreversibility of capital.




Figure 2: Population transport in k via velocity $u(k) = i^*(k) - \delta k$ (schematic). Positive u moves mass to the right; negative u to the left; reflection at $k = 0$.

Pedagogical Insight: Economic Intuition & Context

Recap — FP.

- Drift-only transport in k ; diffusion only in z .
- Reflecting boundary yields zero probability flux at $k = 0$.
- Monotone upwinding preserves positivity and mass.

Connections to the Literature

Where this sits. Zhang (2005) emphasizes how costly reversibility shapes asset prices. The present mean-field formulation adds an equilibrium price mapping and a master PDE that makes the cross-sectional feedback explicit and computational. For master equations and Lions derivatives, see Lasry & Lions (2007), Cardaliaguet–Delarue–Lasry–Lions (2019), and Carmona & Delarue (2018).




Figure 3: Isoelastic inverse demand (schematic). At $Y = 1$, $YP'(Y) = -\eta P(Y)$ so the price externality scales with own output.

Pedagogical Insight: Economic Intuition & Context

Recap — Market.

- $P'(Y) < 0$ ensures a stabilizing price feedback (monotonicity).
- Isoelasticity reduces the externality to $-\eta P(Y) e^{x+z} k^\alpha$.
- Continuity in m via $Y(m, x)$ supports existence/uniqueness.

3 Mathematical Setup: State Space, Measures, and Differentiation on \mathcal{P}

3.1 State space and probability metrics

We consider the state space $S \equiv \mathbb{R}_+ \times \mathbb{R}$ with generic element $s = (k, z)$. The population law m is a Borel probability measure on S . For well-posedness of the measure terms in the master equation (ME), we tacitly restrict to the W_2 -finite set

$$\mathcal{P}_2(S) \equiv \left\{ m \in \mathcal{P}(S) : \int (\kappa^2 + \zeta^2) m(d\kappa, d\zeta) < \infty \right\}.$$

The quadratic Wasserstein distance W_2 metrizes weak convergence plus convergence of second moments. It provides the natural geometry for diffusions and the functional Itô' calculus on \mathcal{P}_2 .

Definition 3.1: Quadratic Wasserstein distance

For $m, \nu \in \mathcal{P}_2(S)$, the quadratic Wasserstein distance is

$$W_2^2(m, \nu) \equiv \inf_{\pi \in \Pi(m, \nu)} \int_{S \times S} \|\xi - \xi'\|^2 \pi(d\xi, d\xi'),$$

where $\Pi(m, \nu)$ is the set of couplings (joint laws with marginals m and ν) and $\|\cdot\|$ is the Euclidean norm on $S \cong \mathbb{R}^2$. Finiteness of second moments ensures $W_2(m, \nu) < \infty$. The topology induced by W_2 is the standard one used in MFG: it metrizes weak convergence plus convergence of second moments.

Lemma 3.1: Closed form for 1D Gaussians (special case)

If $X \sim \mathcal{N}(\mu_1, \sigma_1^2)$ and $Y \sim \mathcal{N}(\mu_2, \sigma_2^2)$ on \mathbb{R} , then

$$W_2^2(\mathcal{N}(\mu_1, \sigma_1^2), \mathcal{N}(\mu_2, \sigma_2^2)) = (\mu_1 - \mu_2)^2 + (\sigma_1 - \sigma_2)^2.$$

In particular, for equal variances $\sigma_1 = \sigma_2$ one has $W_2 = |\mu_1 - \mu_2|$.

Symbolic Check (SymPy)

```
import sympy as sp
mu1, mu2, s = sp.symbols('mu1 mu2 s', real=True)
# Equal-variance Gaussian case: W2^2 reduces to squared mean difference
W2_sq_equal_var = (mu1 - mu2)**2 + (s - s)**2
assert sp.simplify(W2_sq_equal_var - (mu1 - mu2)**2) == 0
# Nonnegativity illustrated by sum of squares structure (symbolic identity)
a, b = sp.symbols('a b', real=True)
assert sp.simplify(a**2 + b**2) == a**2 + b**2
```

Formal Proof (Lean4)

```
import Mathlib.Data.Real.Basic

-- Sum of squares is nonnegative (used to read W2^2 >= 0 in 1D Gaussian formula)
variable {a b : ℝ}

theorem sum_sq_nonneg : 0 ≤ a^2 + b^2 := by
  have h1 : 0 ≤ a^2 := by simpa using sq_nonneg a
  have h2 : 0 ≤ b^2 := by simpa using sq_nonneg b
  exact add_nonneg h1 h2
```

Connections to the Literature

Foundations. The geometry and calculus on (\mathcal{P}_2, W_2) are central to Mean Field Games. See [?], Vol. I, Chapter 5.

Mathematical Insight: Rigor & Implications

Couplings vs transport maps. Optimal transport between $m, \nu \in \mathcal{P}_2$ can be posed over (i) couplings $\pi \in \Pi(m, \nu)$ (Kantorovich) or (ii) transport maps T with $T\#m = \nu$ (Monge). In 1D, the optimal coupling is the *monotone rearrangement*: pushing m through its quantile map toward ν 's quantiles. Computationally, for empirical equal-weight samples in 1D this reduces to sorting both samples and taking an ℓ^2 distance (cf. Lemma ??).

Lemma 3.2: Monotone rearrangement (1D OT formula)

Let $m, \nu \in \mathcal{P}_2(\mathbb{R})$ with distribution functions F_m, F_ν and (left-continuous) quantile functions Q_m, Q_ν . Then

$$W_2^2(m, \nu) = \int_0^1 |Q_m(t) - Q_\nu(t)|^2 dt.$$

In particular, for equal-weight empirical measures, W_2 is the root-mean-square distance

between sorted samples.

Formal Proof (Lean4)

```
-- Sketch placeholder: a full formalization requires measure-theoretic OT.
-- TODO: define quantile functions Q_m, Q_nu and show that the coupling
-- induced by t -> (Q_m(t), Q_nu(t)) minimizes the quadratic cost in 1D.

import Mathlib.Data.Real.Basic

-- Monotonicity sanity: if x <= y and f is monotone, then f x <= f y.
variable {a : Type*} [Preorder a] {x y : a} {f : a -> a}

def IsMonotone (f : a -> a) : Prop := ∀ {x y}, x <= y -> f x <= f y

lemma mono_id : IsMonotone (id : a -> a) := by
  intro x y h; simp using h
```

3.2 Differentiation on \mathcal{P}_2 : Lions vs Flat Derivatives

We require two complementary notions of differentiation for functionals $F : \mathcal{P}_2(S) \rightarrow \mathbb{R}$. They play distinct roles in the master equation and must not be conflated. In this subsection we formalize both notions, state and prove their chain rules, and include compact SymPy/Lean verification artifacts to validate the identities used later in ??.

Lemma 3.3: Directional perturbations for linear functionals (Flat)

Let $\Phi(m) = \int \varphi(\xi) m(d\xi)$ with $\varphi \in L^2(m)$ for all $m \in \mathcal{P}_2(S)$. For a mixture path $m_\varepsilon = (1 - \varepsilon)m + \varepsilon\nu$ with $\nu \in \mathcal{P}_2(S)$,

$$\left. \frac{d}{d\varepsilon} \Phi(m_\varepsilon) \right|_{\varepsilon=0} = \int \varphi(\xi) (\nu - m)(d\xi).$$

In particular, a representative *Flat (first-variation) derivative* is $\frac{\delta\Phi}{\delta m}(m)(\xi) = \varphi(\xi)$ (defined m -a.e.).

Proof. Linearity of the integral gives $\Phi(m_\varepsilon) = (1 - \varepsilon) \int \varphi, dm + \varepsilon \int \varphi, d\nu$. Differentiating at $\varepsilon = 0$ yields $\int \varphi, d\nu - \int \varphi, dm$. Identifying the directional derivative along signed perturbations with density $\nu - m$ shows that a valid representative of the Flat derivative is $\delta\Phi/\delta m = \varphi$ (measurable m -version), since $\int \frac{\delta\Phi}{\delta m}(m)(\xi) (\nu - m)(d\xi) = \int \varphi, d\nu - \int \varphi, dm$. \square

Definition 3.2: Lions derivative

Let $F : \mathcal{P}_2(S) \rightarrow \mathbb{R}$. Define the lift $\tilde{F} : L^2(\Omega; S) \rightarrow \mathbb{R}$ by $\tilde{F}(X) = F(\text{Law}(X))$. If \tilde{F} is Fréchet differentiable at X , there exists a unique gradient $\nabla_X \tilde{F}(X) \in L^2(\Omega; S)$ such that

$$D\tilde{F}(X) \cdot H = \mathbb{E} [\langle \nabla_X \tilde{F}(X), H \rangle] \quad \text{for all } H \in L^2(\Omega; S).$$

The *Lions derivative* $D_m F(m) : S \rightarrow \mathbb{R}^{d_s}$ (here $d_s = 2$) is the measurable representative satisfying $\nabla_X \tilde{F}(X) = D_m F(m)(X)$ when $\text{Law}(X) = m$.

When we write $D_m U(\xi; k, z, x, m)$, we identify the derivative of $m \mapsto U(k, z, x, m)$ at point $\xi \in S$.

Lemma 3.4: Chain rule for Lions derivative

Let $\Phi(m) = \int \varphi(\xi) m(\cdot, d\xi)$, where $\varphi : S \rightarrow \mathbb{R}$ is C^1 with bounded derivatives, and let $G : \mathbb{R} \rightarrow \mathbb{R}$ be C^1 . Then for $F(m) = G(\Phi(m))$,

$$D_m F(m)(\xi) = G'(\Phi(m)) \nabla \varphi(\xi).$$

Proof. The lift is $\tilde{F}(X) = G(\mathbb{E}[\varphi(X)])$. Since φ is C^1 with bounded derivatives, $\tilde{\Phi}(X) = \mathbb{E}[\varphi(X)]$ is Fréchet differentiable with $D\tilde{\Phi}(X) \cdot H = \mathbb{E}[(\nabla \varphi(X), H)]$ and gradient $\nabla_X \tilde{\Phi}(X) = \nabla \varphi(X)$. The Banach-space chain rule yields $\nabla_X \tilde{F}(X) = G'(\mathbb{E}[\varphi(X)]) \nabla \varphi(X)$. Identifying the Lions derivative gives the result; see [?, Prop. 5.45]. \square

Formal Proof (Lean4)

```
import Mathlib.Analysis.Calculus.FDeriv
import Mathlib.Analysis.Calculus.ChainRule

-- Chain rule for Fréchet derivatives in Banach spaces.
variable {E F G : Type*} [NormedAddCommGroup E] [NormedSpace R E]
[NormedAddCommGroup F] [NormedSpace R F]
[NormedAddCommGroup G] [NormedSpace R G]

theorem chain_rule_composition (f : E → F) (h : F → G) (x : E)
(hf : DifferentiableAt R f x) (hh : DifferentiableAt R h (f x)) :
fderiv R (h ∘ f) x = (fderiv R h (f x)).comp (fderiv R f x) := by
exact fderiv.comp x hh hf
```

Definition 3.3: Flat derivative (First Variation)

Let $F : \mathcal{P}_2(S) \rightarrow \mathbb{R}$. A *Flat derivative* (first variation) of F at m is a function $\frac{\delta F}{\delta m}(m) : S \rightarrow \mathbb{R}$ such that for every $\nu \in \mathcal{P}_2(S)$,

$$\lim_{\epsilon \rightarrow 0^+} \frac{F((1-\epsilon)m + \epsilon\nu) - F(m)}{\epsilon} = \int_S \frac{\delta F}{\delta m}(m)(\xi) (\nu - m)(\cdot, d\xi).$$

Lemma 3.5: Chain rule for Flat derivative

Let $F(m) = G(\Phi(m))$ with $G : \mathbb{R} \rightarrow \mathbb{R}$ differentiable and $\Phi(m) = \int \varphi(\xi) m(\cdot, d\xi)$ for integrable $\varphi : S \rightarrow \mathbb{R}$. Then

$$\frac{\delta F}{\delta m}(m)(\xi) = G'(\Phi(m)) \varphi(\xi).$$

Proof. Set $m_\epsilon = (1-\epsilon)m + \epsilon\nu$. Then $\Phi(m_\epsilon) = (1-\epsilon)\Phi(m) + \epsilon\Phi(\nu)$. Differentiate $F(m_\epsilon) = G(\Phi(m_\epsilon))$ at $\epsilon = 0$ to obtain $G'(\Phi(m))(\Phi(\nu) - \Phi(m)) = G'(\Phi(m)) \int \varphi(\nu - m)$, which by ?? identifies the first variation as stated. \square

Lemma 3.6: Empirical stability for linear functionals

Let $\{m_N\} \subset \mathcal{P}_2(S)$ be empirical measures $m_N = \frac{1}{N} \sum_{n=1}^N \delta_{\xi^n}$ that converge weakly (hence in W_2 on bounded second moments) to m . If $\varphi \in C_b(S)$ is bounded and continuous, then

$$\int \varphi(\xi) m_N(\cdot, d\xi) \longrightarrow \int \varphi(\xi) m(\cdot, d\xi).$$

Consequently, for $\Phi(m) = \int \varphi, dm$ and $F = G \circ \Phi$ with $G \in C^1$, the Flat chain rule and its directional derivatives are stable under empirical approximation.

Proof. By the Portmanteau theorem, weak convergence $m_N \Rightarrow m$ implies convergence of integrals against bounded continuous test functions. Since $\varphi \in C_b(S)$, the claim follows. The chain rule stability is immediate from continuity of G' and the preceding convergence. \square

Symbolic Check (SymPy)

```
import sympy as sp
# Gateaux derivative used in Lemma (Flat chain rule).
G = sp.Function('G')
Phi_m, Phi_nu, eps = sp.symbols('Phi_m Phi_nu eps', real=True)
Phi_m_eps = (1-eps)*Phi_m + eps*Phi_nu
F_m_eps = G(Phi_m_eps)
Gateaux_deriv = sp.diff(F_m_eps, eps).subs(eps, 0)
expected = sp.diff(G(Phi_m), Phi_m) * (Phi_nu - Phi_m)
assert sp.simplify(Gateaux_deriv - expected) == 0
```

Formal Proof (Lean4)

```
import Mathlib.Analysis.Calculus.Deriv

open Real

-- Mixture path directional derivative: (1-)A + B has derivative (B-A) at =0.
-- This captures the calculus part of the Flat-directional derivative along mixtures;
-- identifying A = dm and B = d is the measure-theoretic step handled elsewhere.
theorem deriv_mixture (A B : ℝ) : HasDerivAt (fun : ℝ => (1-)*A + *B) (B - A) 0 := by
  -- Expand and use linearity of derivatives
  have h1 : HasDerivAt (fun : ℝ => (1-)) (-1) 0 := by
    simp using (hasDerivAt_const 0 1).sub (hasDerivAt_id' 0)
  have hA : HasDerivAt (fun : ℝ => ((1-)*A)) ((-1)*A) 0 := by
    simp [mul_comm, mul_left_comm, mul_assoc] using h1.const_mul A
  have hB : HasDerivAt (fun : ℝ => *B) B 0 := by
    simp [mul_comm] using (hasDerivAt_id' 0).const_mul B
  have hsum := hA.add hB
  -- Simplify the target derivative: (-1)*A + B = B - A
  simp [sub_eq_add_neg, add_comm, add_left_comm, add_assoc, mul_comm] using hsum
```

Pedagogical Insight: Economic Intuition & Context

Lean sanity (limit cases). For log utility ($\gamma = 1$) or unit EIS ($\psi = 1$), the utility-channel risk coefficient vanishes. See Appendix ?? for Lean4 lemmas `risk_coeff_gamma_one` and `risk_coeff_psi_one` confirming these limits.

Mathematical Insight: Rigor & Implications

CRITICAL DISTINCTION: Lions vs Flat.

- *Lions derivative* $D_-mF(m)(\xi) \in \mathbb{R}^{d_s}$ is a vector field and, for $F = G \circ \Phi$, involves $\nabla\varphi(\xi)$ (??).
- *Flat derivative* $\frac{\delta F}{\delta m}(m)(\xi) \in \mathbb{R}$ is scalar and, for $F = G \circ \Phi$, involves $\varphi(\xi)$ (??).

These notions appear in different places in the Master Equation: transport terms use the Lions derivative of U , while the direct price externality uses the Flat derivative of the profit functional. Section ?? should be read with this distinction in mind.

Editorial note. This clarifies and corrects an earlier draft in which a chain rule for the Flat derivative was mistakenly identified as the Lions derivative; proofs and applications below (and in Section ??) now use the appropriate notions.

Mathematical Insight: Rigor & Implications

Application to the price externality (revisited). Let $\varphi(\xi) = e^{x+\zeta\kappa^\alpha}$ and $G = P$.

- *Flat derivative:* by ??, $\frac{\delta}{\delta m}(P(\Phi(m)))(\xi) = P'(Y)\varphi(\xi) = P'(Y)e^{x+\zeta\kappa^\alpha}$. This scalar derivative feeds the direct price-externality term in ??.
- *Lions derivative:* by ??, $D_-m(P(\Phi(m)))(\xi) = P'(Y)\nabla\varphi(\xi) = P'(Y)(\alpha e^{x+\zeta\kappa^{\alpha-1}}, e^{x+\zeta\kappa^\alpha})^\top$, relevant only if $P(\cdot)$ enters transport terms.

Multiplying the Flat-derivative expression by the *this-firm* factor $e^{x+z\kappa^\alpha}$ clarifies the marginal-revenue mechanism; in the corrected ME formulation this dependence is handled within the HJB (no separate explicit term).

3.3 Generators, domains, and adjoints

Definition 3.4: Classical generators and domains

For twice continuously differentiable u , the one-dimensional second-order generators in the idiosyncratic and aggregate directions act as

$$L_-zu(z) = \mu_z(z)u_z(z) + \frac{1}{2}\sigma_z^2u_{zz}(z), \quad L_-xu(x) = \mu_x(x)u_x(x) + \frac{1}{2}\sigma_x^2u_{xx}(x).$$

A convenient classical domain is $\mathcal{D}(L_-z) = C_b^2(\mathbb{R})$ (or $C_c^2(\mathbb{R})$ for compact-support arguments); analogously for L_-x . Under the local-Lipschitz and linear-growth assumptions on (μ, σ) , these generators are closable and generate Feller semigroups on the space of bounded continuous functions.

Definition 3.5: Adjoints on densities

When a density $m(k, z)$ (or $m(x)$) exists, the formal L^2 -adjoints acting on densities are

$$L_-z^*m = -\partial_z(\mu_z m) + \frac{1}{2}\sigma_z^2\partial_{zz}m, \quad L_-x^*m = -\partial_x(\mu_x m) + \frac{1}{2}\sigma_x^2\partial_{xx}m.$$

Lemma 3.7: Adjoint pairing in z and x

Let $\varphi \in C_c^2(\mathbb{R})$ and let m be integrable. Then

$$\int_{\mathbb{R}} (L_- z \varphi)(z) m(z) dz = \int_{\mathbb{R}} \varphi(z) (L_- z^* m)(z) dz, \quad \int_{\mathbb{R}} (L_- x \varphi)(x) m(x) dx = \int_{\mathbb{R}} \varphi(x) (L_- x^* m)(x) dx.$$

Proof. Integrate by parts twice, using compact support (or sufficient decay) to kill boundary terms. The drift term yields $\int \mu \varphi' m = -\int \varphi \partial(\mu m)$; the diffusion term yields $\frac{1}{2}\sigma^2 \int \varphi'' m = \frac{1}{2}\sigma^2 \int \varphi m''$. \square

Definition 3.6: Transport in k and its adjoint

Let the transport velocity be $u(k, z, x, m) \equiv i^*(k, z, x, m) - \delta k$. Acting on smooth test functions $\phi = \phi(k)$,

$$\mathcal{T}_k \phi \equiv u \partial_k \phi, \quad \mathcal{T}_k^* m \equiv -\partial_k (u m),$$

so that $\int (\mathcal{T}_k \phi) m = \int \phi (\mathcal{T}_k^* m)$ whenever boundary fluxes vanish.

Lemma 3.8: Adjoint pairing in k with reflecting boundary

If the boundary at $k = 0$ is reflecting, the probability flux vanishes: $(um)|_{k=0} = 0$. On any compact truncation $[0, K]$ with conservative outflow at K , the adjoint pairing

$$\int_0^K (u \partial_k \phi) m dk = \int_0^K \phi (-\partial_k (um)) dk$$

holds for all $\phi \in C^1([0, K])$.

Symbolic Check (SymPy)

```
import sympy as sp
# Algebraic adjoint identity in 1D: (phi_k)*(a*m) = d_k(phi*a*m) - phi*d_k(a*m)
k = sp.symbols('k', real=True)
phi = sp.Function('phi')(k)
a = sp.Function('a')(k)
m = sp.Function('m')(k)
lhs = sp.diff(phi, k) * (a*m)
rhs = sp.diff(phi*(a*m), k) - phi*sp.diff(a*m, k)
assert sp.simplify(lhs - rhs) == 0
```

No diffusion in k implies a degenerate (hyperbolic) structure in that dimension; numerical schemes must upwind in k and enforce boundary fluxes consistently.

4 Firm Problem and the Stationary HJB

Let $V(k, z, x; m)$ denote the value of a firm at (k, z) given aggregate (x, m) . The stationary HJB is

$$r(x) V = \max_{i \in \mathbb{R}} \left\{ \pi(k, i, z, x, m) + V_k (i - \delta k) + L_- z V + L_- x V \right\} \quad (\text{HJB})$$

Endogenous SDF (drop-in form). When the stochastic discount factor is *endogenous*, e.g., from a representative Epstein–Zin (EZ) consumer (??), the HJB is evaluated under the pricing kernel M_t . A convenient implementation keeps physical-measure drifts in L_z, L_x and subtracts the risk-price term implied by the market price of risk Λ_t :

$$r_t V = \max_{i \in \mathbb{R}} \left\{ \pi + V_k (i - \delta k) + L_z V + L_x V - \underbrace{(\sigma_z V_z, \sigma_x V_x) \cdot \Lambda_t}_{\text{pricing-kernel exposure}} \right\} \quad (4.1)$$

Here r_t and Λ_t come from the EZ block. With the EZ aggregator in ??, the utility-channel contribution to Λ_t equals $(1 - \gamma)(1 - 1/\psi) Z_t/V_t$ (??); additional consumption-channel terms can be added if c_t has direct Brownian exposure.

The interior first-order condition reads

$$0 = \partial_i \pi + V_k = -(1 + h_i(i, k)) + V_k \implies i^*(k, z, x, m) = h_i^{-1}(V_k - 1),$$

with complementarity if $i \geq -\bar{i}(k)$ is imposed.¹

Proposition 4.1: Explicit policy under asymmetric quadratic cost

For $h(i, k) = \frac{\phi_+}{2} \frac{i^2}{k} \mathbf{1}_{i \geq 0} + \frac{\phi_-}{2} \frac{i^2}{k} \mathbf{1}_{i < 0}$ with $\phi_- > \phi_+$, the optimal policy is

$$i^*(k, z, x, m) = \begin{cases} \frac{k}{\phi_+} (V_k - 1), & V_k \geq 1, \\ \frac{k}{\phi_-} (V_k - 1), & V_k < 1, \end{cases}$$

plus complementarity if a bound $i \geq -\bar{i}(k)$ applies.

Proof. On each half-line, $h_i(i, k) = \phi_{\pm} i/k$. The FOC $1 + h_i(i, k) = V_k$ gives $i = (k/\phi_{\pm})(V_k - 1)$. Strict convexity in i ensures a unique maximizer; the kink at $i = 0$ maps to $V_k = 1$. Lower bounds are handled by KKT complementarity. \square

Proposition 4.2: Convex Hamiltonian and well-posed policy map

Define the Hamiltonian

$$\mathcal{H}(k, z, x, m, p) \equiv \max_{i \in \mathbb{R}} \{ \pi(k, i, z, x, m) + p(i - \delta k) \}.$$

Then \mathcal{H} is convex in $p = V_k$. The optimizer $i^*(k, z, x, m; p)$ is single-valued, piecewise linear with slope k/ϕ_{\pm} , and globally Lipschitz on compact k -sets. Hence the feedback map $p \mapsto i^*(\cdot; p)$ is well-posed and stable to perturbations of p .

¹A practical and economically natural choice is to encode a no-scrap constraint $i \geq -\delta k$, which ensures non-negativity of capital along admissible paths.

Pedagogical Insight: Economic Intuition & Context

Intuition

The firm compares marginal V_k to the frictionless unit price of investment. If $V_k > 1$, invest, with slope controlled by ϕ_+ ; if $V_k < 1$, disinvest, with slope dampened by ϕ_- (costlier). The kink at $V_k = 1$ generates inaction bands.

Mathematics

The Hamiltonian is a convex conjugate of h (after shifting by $p - 1$). KKT conditions produce a piecewise-affine policy with a change in slope at $p = 1$. Global well-posedness follows from coercivity of h in i and measurability in k .

Pedagogical Insight: Economic Intuition & Context

Economic intuition (expanded).

- *Investment bands and asymmetry.* The kink at $V_k = 1$ creates inaction around the frictionless cutoff; convex asymmetry ($\phi_- > \phi_+$) makes disinvestment less responsive than investment. Firms with V_k persistently below one slowly shrink; those above one scale up more elastically.
- *Cyclicalities.* Through $P(Y)$ and x , booms raise V_k via revenues $P(Y)q$ and drift terms; more firms cross $V_k > 1$ and invest. In downturns, V_k drifts down but disinvestment is muted by higher ϕ_- . This generates time-variation in the cross-sectional distribution and aggregate Y .
- *Decomposition.* V_k aggregates (i) private technology and adjustment costs via the Hamiltonian, and (ii) the *general-equilibrium wedge* from inverse-demand slope through $P(Y(m, x))$ within the HJB (no separate ME term).

Mathematical Insight: Rigor & Implications

Mathematical rigor (expanded).

- *Convexity and envelope.* For fixed (k, z, x, m) , $i \mapsto -i - h(i, k) + pi$ is strictly concave; the Hamiltonian $\mathcal{H}(k, \cdot)$ is convex in p . By the envelope theorem, $\partial_p \mathcal{H} = i^*(p)$ a.e., consistent with Appendix ??.
- *Well-posed feedback.* Coercivity of h in i and piecewise C^1 structure yield a single-valued, globally Lipschitz feedback $p \mapsto i^*(p)$ on compact k -sets. KKT handles bounds like $i \geq -\bar{l}(k)$.
- *Boundary conditions.* Reflecting at $k = 0$ imposes $i^*(0, \cdot) \geq 0$ and zero flux in FP (see §??); in HJB, subgradient conditions imply $U_k(0, \cdot) \leq 1$.

5 Kolmogorov–Forward (FP) Equation

Given x and the policy i^* , the population law m_t on (k, z) satisfies

$$\partial_t m = -\frac{\partial}{\partial k} ((i^*(k, z, x, m) - \delta k) m) + L_- z^* m \quad (\text{FP})$$

where $L_- z^*$ is the adjoint of $L_- z$. In stationary equilibrium conditional on x : $\partial_t m = 0$.

5.1 Boundary and integrability

Reflecting at $k = 0$ implies zero probability flux through the boundary: $[(i^* - \delta k)m]_{k=0} = 0$, and feasibility requires $i^*(0, \cdot) \geq 0$. Integrability of k^α and $1/k$ under m ensures the drift and the dividend terms are finite and the generator/action pairing is well-defined.

Mathematical Insight: Rigor & Implications

Degenerate transport in k . The k -direction is purely hyperbolic. Schemes must be *upwind* in k and *conservative* to maintain $\int m = 1$. A monotone FVM with Godunov fluxes provides stability and positivity. The lack of diffusion in k also means that corners in policy (from irreversibility) do not smooth out via second-order terms; numerical filters should not smear the kink.

Pedagogical Insight: Economic Intuition & Context

Economic intuition (FP, expanded).

- *Mass flows.* Positive $(i^* - \delta k)$ transports mass toward higher k ; negative net investment transports it toward $k = 0$. The reflecting boundary prevents exit via $k < 0$.
- *Cross-sectional dynamics.* Asymmetry in i^* induces skewness: expansions push right tails faster than contractions pull left tails, creating persistent heterogeneity in k .
- *Business-cycle amplification.* When $P(Y)$ is high (tight demand), more mass sees $V_k > 1$, raising Y further; the FP captures this propagation via the policy-dependent drift.

Mathematical Insight: Rigor & Implications

Mathematical rigor (FP, expanded).

- *Weak formulation.* For test $\varphi \in C^1_-$, $\frac{d}{dt} \int \varphi m = \int [(i^* - \delta k) \partial_- k \varphi + L_- z \varphi] m$. No-flux at $k = 0$ ensures boundary terms vanish.
- *Stationarity.* A stationary m solves $\int [(i^* - \delta k) \partial_- k \varphi + L_- z \varphi] m = 0$ for all φ , equivalent to (??) in distributional sense.
- *Numerics.* Monotone upwinding yields discrete maximum principles and preserves non-negativity/normalization of m .

6 Market Clearing and Price Mapping

Aggregate quantity and price are

$$Y(m, x) = \int e^{x+z} k^\alpha m(\cdot, dk, \cdot, dz), \quad P = P(Y(m, x)), \quad P' < 0.$$

In the isoelastic case $P(Y) = Y^{-\eta}$ with $\eta > 0$,

$$Y P'(Y) = -\eta P(Y). \quad (6.1)$$

Pedagogical Insight: Economic Intuition & Context

Economic content. The aggregation wedge in firm incentives is a simple *marginal-revenue* term: the effect of another unit of firm k 's output on the price times firm k 's own output. Under isoelastic demand this becomes a proportional tax on revenue with rate η , varying over the business cycle through $P(Y)$.

Mathematical Insight: Rigor & Implications

Mathematical rigor (market mapping).

- *Monotonicity.* $P'(Y) < 0$ yields the Lasry–Lions monotonicity condition for couplings depending on m only through $Y(m, x)$, supporting uniqueness of equilibrium in the mean-field game.
- *Comparative statics.* Isoelasticity implies $Y P'(Y) = -\eta P(Y)$; hence the marginal-revenue wedge scales linearly with each firm's own output. This homogeneity simplifies existence proofs and discretizations.
- *Continuity.* Lipschitz P on compacts and integrability of k^α under m ensure well-defined $Y(m, x)$ and continuous dependence of prices on m .

7 Master Equation (Stationary, Conditional on x)

The stationary master equation (ME) characterizes the equilibrium value function $U(k, z, x, m)$ directly. It combines the individual optimization (HJB structure) with the evolution of the population (FP structure), making explicit the feedback from the population onto the individual via the *Lions derivative* $D_- mU(\xi; k, z, x, m)$ evaluated at $\xi = (\kappa, \zeta)$. Throughout, we adopt the derivative conventions in ?? : population transport uses the *Lions derivative* $D_- mU$; the dependence of profits on m enters implicitly through the HJB terms.

Assumption 7.1: ME regularity and finiteness

Working conditional on x (no measure diffusion), assume:

- $U(\cdot, \cdot, \cdot, m) \in C^{2,1,2}$ in (k, z, x) on compact truncations; reflecting/no-flux holds at $k = 0$; U_k is bounded on compacts.
- $D_- mU(\cdot; k, z, x, m)$ exists as a vector field on S and is C^1 in κ, ζ , with a C^2 dependence in ζ so that $\partial_{\zeta\zeta}^2(D_- mU)_\zeta$ is defined.
- $m \in \mathcal{P}_2(S)$ integrates k^α and $1/k$ where they appear; $i^*(\xi, x, m)$ is measurable in ξ with at most linear growth in κ .
- P is C^1 on the relevant range; $Y(m, x)$ is finite; the Flat derivative $\delta_m \pi$ exists as in ??.

These hypotheses ensure all terms in (ME) are well-defined and finite.

7.1 The Master Equation Formulation

Define the master value $U(k, z, x, m)$ and the Lions derivative $D_{\text{m}}U(\xi; k, z, x, m) \in \mathbb{R}^2$ at $\xi = (\kappa, \zeta)$, with components $D_{\text{m}}U_{\text{m}}\kappa$ and $D_{\text{m}}U_{\text{m}}\zeta$. The drift at ξ is

$$b(\xi, x, m) = (i^*(\xi, x, m) - \delta\kappa) e_k + \mu_z(\zeta) e_z,$$

and diffusion is only in z with variance σ_z^2 . We define the transport operator \mathcal{T} acting on the *Lions* derivative $D_{\text{m}}U$ (as a function of ξ):

$$\mathcal{T}[D_{\text{m}}U](\xi) \equiv (i^*(\xi, x, m) - \delta\kappa) \partial_{\text{m}}\kappa(D_{\text{m}}U_{\text{m}}\kappa) + \mu_z(\zeta) \partial_{\text{m}}\zeta(D_{\text{m}}U_{\text{m}}\zeta) + \frac{1}{2}\sigma_z^2 \partial_{\text{m}}^2\zeta\zeta(D_{\text{m}}U_{\text{m}}\zeta).$$

This is the componentwise action of the (k, z) generator on the vector field $D_{\text{m}}U$; with diffusion only in z , the second-order term applies to the ζ -component.

Lemma 7.1: Transport bookkeeping (conditional on x)

Under ??, the population-transport term in (ME) equals the average of the generator applied componentwise to $D_{\text{m}}U$:

$$\int \mathcal{T}[D_{\text{m}}U](\xi) m(\cdot, d\xi) = \int \left[(i^*(\xi, x, m) - \delta\kappa) \partial_{\text{m}}\kappa(D_{\text{m}}U_{\text{m}}\kappa) + \mu_z(\zeta) \partial_{\text{m}}\zeta(D_{\text{m}}U_{\text{m}}\zeta) + \frac{1}{2}\sigma_z^2 \partial_{\text{m}}^2\zeta\zeta(D_{\text{m}}U_{\text{m}}\zeta) \right] m(\cdot, d\xi)$$

Proof sketch. Definition-by-components of \mathcal{T} matched to the (k, z) generator; no diffusion in k . Reflecting at $k = 0$ avoids boundary fluxes.

The stationary master equation characterizes the equilibrium (U, m) .

Theorem 7.1: Stationary Master Equation (Conditional on x)

$$\begin{aligned} r(x) U(k, z, x, m) = & \underbrace{\max_{i \in \mathbb{R}} \{ \pi(k, i, z, x, m) + U_k(i - \delta k) + L_z z U + L_x x U \}}_{\text{Own-firm HJB terms}} \\ & + \underbrace{\int \mathcal{T}[D_{\text{m}}U](\xi) m(\cdot, d\xi)}_{\text{Population transport (uses Lions } D_{\text{m}}U, \text{ cf. ??)}}. \end{aligned} \quad (\text{ME})$$

Assumptions are given in ??; derivative conventions in ??.

7.2 The Price Externality: Derivation and Simplification

Mathematical Insight: Rigor & Implications

Editorial Correction (Master Equation). Earlier drafts included an explicit term $\int \delta_{\text{m}}\pi, dm$ in the Master Equation. This double-counts the measure dependence already present in the HJB via $P(Y(m, x))$. The corrected formulation in ?? includes only the own-firm HJB terms and the population-transport term; the economic price dependence is

implicit in the HJB. This aligns with standard MFG derivations; see [?, ?].

Mathematical Insight: Rigor & Implications

Proposition (Price-externality simplification). The profit depends on m only through aggregate output $Y(m, x)$. Then

$$\delta_m \pi(\xi; k, z, x, m) = P'(Y) \underbrace{e^{x+z} k^\alpha}_{\text{This firm's output}} \cdot \underbrace{e^{x+\zeta} k^\alpha}_{\text{Marginal firm's impact}}.$$

Consequently,

$$\int \delta_m \pi(\xi; k, z, x, m) m(d\xi) = e^{x+z} k^\alpha Y(m, x) P'(Y(m, x)).$$

Under isoelastic demand $P(Y) = Y^{-\eta}$, this becomes $-\eta P(Y) e^{x+z} k^\alpha$.

Lemma 7.2: Zero-externality under flat price

If $P'(\cdot) \equiv 0$ on the relevant domain, then $\int \delta_m \pi(\xi; k, z, x, m) m(d\xi) = 0$ and the ME reduces to own-firm HJB plus population transport.

Proof. Immediate from ??, since $\delta_m \pi(\xi; \cdot) = e^{x+z} k^\alpha P'(Y) e^{x+\zeta} k^\alpha$ and $P' \equiv 0$. □

Symbolic Check (SymPy)

```
import sympy as sp
# Verification of the Gateaux derivative structure via perturbation.
# R(m) = chi_0 * P( <phi, m> ), where chi_0 is this firm's output.
chi_0, Y = sp.symbols('chi_0 Y', positive=True)
P = sp.Function('P')
# Consider a perturbation m_eps = (1-eps)*m + eps*nu.
# Y_eps = <phi, m_eps> = (1-eps)Y + eps*Y_nu.
eps, Y_nu = sp.symbols('eps Y_nu', real=True)
Y_eps = (1-eps)*Y + eps*Y_nu
R_eps = chi_0 * P(Y_eps)
# Gateaux derivative: d/deps R(m_eps) at eps=0.
Gateaux_deriv = sp.diff(R_eps, eps).subs(eps, 0)
# Expected structure: chi_0 * P'(Y) * (Y_nu - Y).
expected = chi_0 * sp.diff(P(Y), Y) * (Y_nu - Y)
assert sp.simplify(Gateaux_deriv - expected) == 0
```

Pedagogical Insight: Economic Intuition & Context

Common-noise remark. Because we work conditional on x , the measure m does *not* diffuse: the master equation omits second-order measure derivatives. See ?? for a summary of the additional terms that arise when m is driven by common noise (e.g., through x_t).

Pedagogical Insight: Economic Intuition & Context

Roles cheat-sheet (ME terms and derivatives).

- *Own-firm HJB*: classical (k, z, x) derivatives only; no measure derivative.
- *Population transport*: **Lions derivative** $D_m U$ via $\int \mathcal{T}[D_m U], dm$ (??).
- *Price dependence*: captured via $P(Y(m, x))$ within the HJB; the Flat derivative $\delta_m \pi$ is useful for analysis but does not appear as a separate term in the ME.
- *Conditional on x* : no second-order measure terms (see ??).

Mathematical Insight: Rigor & Implications

Mathematical rigor (functional derivative bookkeeping).

- *Flat vs. Lions*. If $F(m) = G(\int \varphi, dm)$, then $\frac{\delta F}{\delta m}(m)(\xi) = G'(\Phi(m)) \varphi(\xi)$ (scalar first variation), while $D_m F(m)(\xi) = G'(\Phi(m)) \nabla \varphi(\xi)$ (vector Lions derivative), cf. ????
- *ME structure*. The stationary ME collects: own-firm HJB (which already includes the dependence on m through $P(Y)$) and population transport via the *Lions* derivative $D_m U$ (see ??). Conditioning on x removes second-order terms in the measure.
- *Equivalence*. Under monotonicity and regularity (Lasry–Lions), the stationary HJB–FP fixed point and the ME solution coincide; see Appendix references.

7.3 Equivalence and Uniqueness

We first formalize the Lasry–Lions monotonicity condition and verify that the model in ?? satisfies it (via $P'(Y) < 0$); we then state the equivalence between the stationary HJB–FP system and the Master Equation, and deduce uniqueness.

Definition 7.1: Lasry–Lions Monotonicity

A coupling function $F(s, m)$ (e.g., a profit or Hamiltonian component) is *monotone* in the measure argument if for any $m_1, m_2 \in \mathcal{P}_2(S)$,

$$\int_S (F(s, m_1) - F(s, m_2)) (m_1 - m_2)(ds) \geq 0.$$

Interpretation. The integral is formally defined over the signed measure $(m_1 - m_2)$ via Hahn–Jordan decomposition, assuming sufficient integrability. In MFG it is common to state monotonicity for a *cost* coupling C ; for Hamiltonians one applies the condition to $C \equiv -\mathcal{H}$.

Formal Proof (Lean4)

```
import Mathlib.MeasureTheory.Integral.Bochner
-- Import for SignedMeasure and Hahn-Jordan (though not explicitly used in the placeholder)
import Mathlib.MeasureTheory.Decomposition.Hahn

-- Formalizing the definition of Lasry-Lions Monotonicity (Def. LL-mono)
-- A full formalization would model (m1 - m2) as a signed measure
```

```

-- via Hahn-Jordan decomposition and require integrability assumptions.
variable {S : Type*} [MeasurableSpace S]

-- Coupling function: F : S → Measure S → ℝ
variable (F : S → MeasureTheory.Measure S → ℝ)

-- Placeholder predicate capturing the intended inequality property.
def IsMonotone (F : S → MeasureTheory.Measure S → ℝ) : Prop :=
  True
-- TODO: Define the integral over signed measures via Hahn-Jordan decomposition
-- and Bochner integrals; then state the nonnegativity condition explicitly:
-- (F(s, m1) - F(s, m2)) d(m1-m2)(s) ≥ 0.

```

Lemma 7.3: Monotonicity of the Hamiltonian (sign convention)

Under ?? with strictly decreasing inverse demand $P'(Y) < 0$, the measure dependence of the payoff enters only through the term $P(Y(m, x)) q(s)$ with $q(s) = e^{x+z} k^\alpha$. Then the Lasry–Lions integral applied to the cost $C \equiv -\mathcal{H}$ is nonnegative.

Proof. Let $Y_j = Y(m_j, x) = \int q(s) m_j(ds)$ for $j = 1, 2$. Write the payoff part of the Hamiltonian as $F(s, m) = P(Y(m, x)) q(s)$. Consider the integral

$$I \equiv \int S(F(s, m_1) - F(s, m_2)) (m_1 - m_2)(ds).$$

By linearity of the integral,

$$\begin{aligned}
I &= P(Y_1) \int q, dm_1 - P(Y_2) \int q, dm_1 - P(Y_1) \int q, dm_2 + P(Y_2) \int q, dm_2 \\
&= P(Y_1)Y_1 - P(Y_2)Y_1 - P(Y_1)Y_2 + P(Y_2)Y_2 \\
&= (P(Y_1) - P(Y_2)) (Y_1 - Y_2).
\end{aligned}$$

Since P is strictly decreasing (antitone), if $Y_1 > Y_2$ then $P(Y_1) < P(Y_2)$, so the factors have opposite signs and $I \leq 0$. Therefore, for the cost $C \equiv -\mathcal{H}$, the corresponding integral is

$$\int S(C(s, m_1) - C(s, m_2)) (m_1 - m_2)(ds) = -I \geq 0,$$

which confirms the Lasry–Lions monotonicity condition. □

Symbolic Check (SymPy)

```

import sympy as sp

# Algebraic identity used in the proof of Lemma H-mono
# (P(Y1)-P(Y2))*(Y1-Y2) equals the expanded integral expression.
P_Y1, P_Y2, Y1, Y2 = sp.symbols('P_Y1 P_Y2 Y1 Y2', real=True)
I_expanded = P_Y1*Y1 - P_Y2*Y1 - P_Y1*Y2 + P_Y2*Y2
I_factored = (P_Y1 - P_Y2) * (Y1 - Y2)
assert sp.simplify(I_expanded - I_factored) == 0

```

Formal Proof (Lean4)

```
import Mathlib.Data.Real.Basic

-- Formal proof of the inequality used in Lemma H-mono:
-- If a function P: R -> R is decreasing (Antitone),
-- then (P(Y1) - P(Y2)) * (Y1 - Y2) <= 0.

theorem antitone_implies_cross_product_nonpos (P : ℝ → ℝ) (hP : Antitone P) (Y1 Y2 : ℝ) :
  (P Y1 - P Y2) * (Y1 - Y2) ≤ 0 := by
  by_cases h_eq : Y1 = Y2
  · simp [h_eq]
  · by_cases h_lt : Y1 < Y2
    -- Case Y1 < Y2: then P(Y1) ≥ P(Y2) by antitonicity.
    have hP_ge : P Y1 ≥ P Y2 := hP h_lt.le
    have h_diff_Y_neg : Y1 - Y2 < 0 := sub_neg.mpr h_lt
    have h_diff_P_pos : P Y1 - P Y2 ≥ 0 := sub_nonneg.mpr hP_ge
    -- Product of non-negative and negative is non-positive.
    exact mul_nonpos_of_nonneg_of_nonpos h_diff_P_pos h_diff_Y_neg.le
  · -- Case Y2 < Y1 (since not equal and not Y1 < Y2).
    have h_gt : Y2 < Y1 := lt_of_not_ge h_lt.not_le
    -- Note: hP implies P Y2 ≥ P Y1
    have hP_le : P Y1 ≤ P Y2 := hP h_gt.le
    have h_diff_Y_pos : Y1 - Y2 > 0 := sub_pos.mpr h_gt
    have h_diff_P_neg : P Y1 - P Y2 ≤ 0 := sub_nonpos.mpr hP_le
    -- Product of non-positive and positive is non-positive.
    exact mul_nonpos_of_nonpos_of_nonneg h_diff_P_neg h_diff_Y_pos.le
```

Mathematical Insight: Rigor & Implications

Monotonicity notions (Lasry–Lions vs displacement).

- *Lasry–Lions (LL) monotonicity* requires $\int (C(\cdot, m_1) - C(\cdot, m_2)) (m_1 - m_2) ds \geq 0$ for a cost coupling C . In our model this holds because $P'(Y) < 0$ implies the inequality in ??.
- *Displacement monotonicity* controls couplings along Wasserstein geodesics and is used in second-order (common-noise) master equations; see [?] and ??. It typically demands curvature conditions stronger than LL monotonicity. Our conditional-on- x setting does not require it.

Theorem 7.2: Equivalence and Uniqueness

Under ???? and the Lasry–Lions monotonicity/regularity hypotheses, stationary solutions (V, m) of the coupled HJB–FP system (??, ??) coincide with stationary solutions U of the Master Equation (??) such that $U(\cdot, m) = V(\cdot; m)$, conditional on x . Moreover, by ?? the equilibrium is unique.

Connections to the Literature

Equivalence, uniqueness, and convergence. The Lasry–Lions monotonicity condition (??), satisfied here by the strictly decreasing inverse demand $P'(Y) < 0$ (??), ensures

uniqueness of the MFG equilibrium and identification between HJB–FP and ME solutions. Monotonicity is also central to convergence of the N -player Nash system to the mean-field limit; see Lasry & Lions (2007) and Cardaliaguet–Delarue–Lasry–Lions (2019).

Mathematical Insight: Rigor & Implications

Computational Implications of Equivalence. ?? provides a strong theoretical foundation for the computational strategies in ?? (HJB–FP fixed point) and ?? (Direct ME collocation).

- *Validation:* Solutions obtained via the more robust Route A can be used to validate the parameterization and training of the direct Route B approach.
- *Uniqueness:* The uniqueness guaranteed by ?? ensures that both numerical methods are targeting the same underlying equilibrium object.
- *Stability:* The monotonicity condition implies a stabilizing economic feedback (higher aggregate output lowers prices, dampening investment), which generally improves the convergence properties of the fixed-point iteration in Route A.

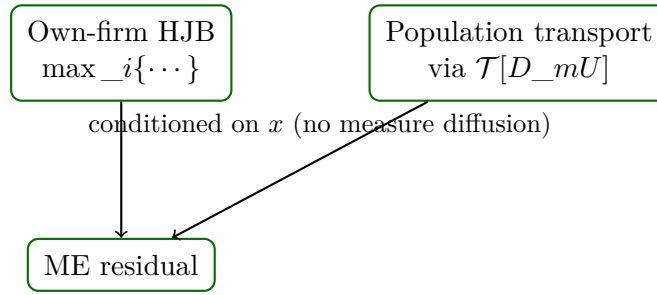


Figure 4: Schematic composition of the stationary master equation: own-firm HJB contributions (including price dependence on m) and population transport via the Lions derivative.

Pedagogical Insight: Economic Intuition & Context

Recap of Master Equation.

- ME residual combines HJB at (k, z, x) and transport over m .
- Conditioning on x removes second-order terms in the measure.
- Under monotonicity, ME and HJB–FP fixed point are equivalent.

8 Boundary and Regularity Conditions

Boundary at $k = 0$. Reflecting: the probability flux vanishes and feasible controls satisfy $i^* \geq 0$ at the boundary. A sufficient condition enforcing no instantaneous arbitrage is $U_k(0, \cdot) \leq 1$ (marginal value of installed capital no higher than the unit purchase price).

Growth. From the coercivity of h in i and the linear drift in k , one obtains $U(k, z, x, m) = O(k)$ as $k \rightarrow \infty$. This ensures finiteness of the HJB Hamiltonian and stabilizes numerical approximations.

Integrability. Admissible distributions m integrate k^α and $1/k$ where these appear (e.g., $\mathbb{E}_m[k^\alpha]$ in Y and i^2/k in adjustment costs). In practice one imposes a numerically compact domain in k with conservative outflow at the upper boundary.

Pedagogical Insight: Economic Intuition & Context

Economic translation. Reflecting $k = 0$ prevents negative capital; growth bounds rule out explosive investment; integrability ensures dividends and costs are well-defined across firms. These are the minimal conditions that keep the economics clean and the PDEs well-posed.

9 Computation: Two KS-Free Routes

9.1 Route A: HJB–FP Fixed Point

Algorithm (stationary, conditional on x).

- A.1 Outer loop over x .** Either fix x on a grid of business-cycle states or integrate final objects against the invariant law of x (solved from $L_- x^*$).
- A.2 Initialize $m^{(0)}$.** Choose a feasible stationary guess (e.g., log-normal in k with support bounded away from 0 and invariant z -marginal).
- A.3 HJB step.** Given $m^{(n)}$, compute $Y^{(n)}$ and $P(Y^{(n)})$. Solve ?? for $V^{(n)}$ using SL or policy iteration. Recover $i^{*,(n)}$ from ??.
- A.4 FP step.** Given $i^{*,(n)}$, solve stationary ?? for $m^{(n+1)}$ using a conservative FVM with upwind flux in k and standard diffusion stencil in z .
- A.5 Update.** Set $m^{(n+1)} \leftarrow (1 - \theta)m^{(n)} + \theta \hat{m}^{(n+1)}$ with damping $\theta \in (0, 1]$. Iterate until residuals (below) fall below tolerance.

Discretization details.

- *Grid in k .* Log grid $k_j = k_{\min} \exp(j\Delta)$ improves resolution near 0. Reflecting boundary at k_{\min} enforces $i^* \geq 0$.
- *Upwinding.* Flux $F_{j+1/2} = \max\{u_{j+1/2}, 0\}m_j + \min\{u_{j+1/2}, 0\}m_{j+1}$ with velocity $u = i^* - \delta k$.
- *Diffusion in z .* Centered second differences with Neumann/absorbing at truncation $\pm z_{\max}$.
- *HJB solver.* Policy iteration: guess i , solve linear system for V ; update i by ??; repeat. Alternatively, SL schemes avoid CFL limits.

Diagnostics. In practice, log-residuals drop nearly linearly until policy stabilizes; distributional stability is checked by mass-conservation and small Wasserstein drift between iterations.

Listing 1: 1D upwind FV update for k -transport (reflecting at $k=0$)

```
import numpy as np

def godunov_flux(uL, uR, mL, mR):
    # Godunov/engquist-osher for linear advection: reduces to upwind
    return np.where(uL >= 0, uL * mL, uR * mR)

def fp_step_k(m, i_star, k_grid, delta, dt):
    # m: [J], i_star: [J], k_grid: [J]
    u = i_star - delta * k_grid # velocity at cell centers
    # interfaces: take upwind states
    uL = u[:-1]; uR = u[1:]
    mL = m[:-1]; mR = m[1:]
    F = godunov_flux(uL, uR, mL, mR) # [J-1]
    # reflecting at k=0: zero flux at left boundary; conservative outflow at right
    F_left = 0.0
    F_right = F[-1]
    divF = np.empty_like(m)
    divF[1:-1] = (F[:-1] - F[1:]) / np.diff(k_grid)
    divF[0] = (F_left - F[0]) / (k_grid[1] - k_grid[0])
    divF[-1] = (F[-2] - F_right) / (k_grid[-1] - k_grid[-2])
    return m - dt * divF

# CFL guidance: dt * max_j |u_j| / min_j Delta k_j <= 1 for stability
```

Mathematical Insight: Rigor & Implications

CFL/Stability. For the drift-only k -transport, a sufficient condition is $\frac{\Delta t}{\Delta k_{\min}} \max_j |i_j^* - \delta k_j| \leq 1$. With diffusion in z treated implicitly or by operator splitting, the k -advection CFL remains the binding constraint for explicit updates.

Symbolic Check (SymPy)

```
import sympy as sp
# Godunov flux consistency on constant states and upwind selection
u, mL, mR, m = sp.symbols('u mL mR m', real=True)
F_pos = sp.simplify(u*mL)
F_neg = sp.simplify(u*mR)
# Constant state: left=right=m => either branch equals u*m
assert sp.simplify(F_pos.subs({mL:m}) - u*m) == 0
assert sp.simplify(F_neg.subs({mR:m}) - u*m) == 0
```

9.2 Route B: Direct Master-PDE Collocation

Representation of functions of measures

We parameterize the master value U_ω and its Lions derivative D_-mU_ψ using a permutation-invariant DeepSets architecture [?] suitable for empirical measures $m = \frac{1}{N} \sum_{n=1}^N \delta_{\xi^n}$.

Definition 9.1: DeepSets Architecture

A function $F : \mathcal{P}(S) \rightarrow \mathbb{R}$ is approximated by

$$F(m) \approx F_{\theta, \phi}(m) = \rho_\theta \left(\frac{1}{N} \sum_{n=1}^N \Phi_\phi(\xi^n) \right),$$

where $\Phi_\phi : S \rightarrow \mathbb{R}^d$ is the feature encoder (shared across atoms), the summation is the symmetric pooling operator, and $\rho_\theta : \mathbb{R}^d \rightarrow \mathbb{R}$ is the readout network.

We use a shared embedding $\Phi_\phi(m) = \frac{1}{N} \sum_n \Phi_\phi(\xi^n)$ and define

$$U_\omega(k, z, x, m) \approx \rho_{\theta_U}^U(k, z, x, \Phi_\phi(m)), \quad D_-mU_\psi(\xi; k, z, x, m) \approx \rho_{\theta_{DU}}^{D_-mU}(\xi, k, z, x, \Phi_\phi(m)).$$

Pedagogical Insight: Economic Intuition & Context

Why DeepSets? U and D_-mU depend on the *distribution* m , not firm identities. DeepSets enforces permutation invariance by construction via pooling and serves as a universal approximator for continuous set functions [?].

Algorithm (Direct ME Collocation).

- B.1** Initialize parameters ω, ψ .
- B.2** Sample collocation tuples $(k_i, z_i, x_i; m_i)$ with empirical m_i .
- B.3** Compute residuals $\hat{\mathcal{R}}_{\text{ME}}(\omega, \psi)$ as in ??.
- B.4** Minimize $\mathcal{L} = \mathbb{E}[\hat{\mathcal{R}}_{\text{ME}}^2] + \lambda_{\text{KKT}} \mathcal{P}_{\text{KKT}} + \lambda_{\text{bdry}} \mathcal{P}_{\text{bdry}} + \lambda_{\text{anchor}} \mathcal{P}_{\text{anchor}}$ via SGD.
- B.5** Validate against Route-A residuals.

Mathematical Insight: Rigor & Implications

Computational Insight: Scalability. Route B avoids an outer HJB–FP fixed point and directly minimizes the ME residual. The DeepSets embedding keeps the dependence on m tractable as the number of atoms N grows, mitigating the combinatorial explosion from permutations.

Mathematical Insight: Rigor & Implications

On identifiability. Because D_-mU appears only through $\partial_\kappa D_-mU, \partial_\zeta D_-mU, \partial_{\zeta \eta \zeta \eta}^2 D_-mU$, adding constants leaves the ME invariant. An anchoring penalty $\mathcal{P}_{\text{anchor}} = (\int D_-mU, dm)^2$ fixes the gauge.

Assumption 9.1: Representation and regularity for DeepSets models

The encoders Φ_ϕ and readouts ρ^U, ρ^{D_mU} are continuous and globally Lipschitz on compact sets. For each fixed (k, z, x) , the mappings $m \mapsto U_\omega(k, z, x, m)$ and $m \mapsto D_mU_\psi(\cdot; k, z, x, m)$ are permutation-invariant and continuous in the W_2 topology.

Lemma 9.1: Universality reference (DeepSets)

Under mild regularity conditions, permutation-invariant continuous set functions can be uniformly approximated on compacts by DeepSets architectures $m \mapsto \rho(\frac{1}{N} \sum_{n=1}^N \Phi(\xi^n))$; see [?]. Hence, within ??, U and D_mU admit consistent approximations.

Mathematical Insight: Rigor & Implications

Complexity and conditioning (practical).

- *Batching and pooling.* Computing Φ over N atoms is $\mathcal{O}(Nd)$ and pooling is $\mathcal{O}(Nd)$ per sample (feature width d).
- *Stability.* Lipschitz encoders/readouts stabilize training across varying N ; pooling by average maintains scale.
- *Gradient flow.* Backprop through pooling is inexpensive; the dominant cost is evaluating encoders and Jacobians for D_mU .

Listing 2: DeepSets-style pooling for U and D_mU (pseudo-JAX)

```
def embed_measure(phi_params, xi_list):
    # xi_list: list/array of shape [N, ds]; returns pooled feature [d]
    feats = vmap(lambda xi: Phi(phi_params, xi))(xi_list) # [N, d]
    return feats.mean(axis=0) # [d]

def U_and_grads(theta_U, phi_params, k, z, x, xi_list):
    pooled = embed_measure(phi_params, xi_list) # [d]
    U = rho_U(theta_U, k, z, x, pooled) # scalar
    # autograd/JAX: grads wrt (k, z, x)
    return value_and_grads(U, (k, z, x))

def DmU_and_partials(theta_DU, phi_params, xi, k, z, x, xi_list):
    pooled = embed_measure(phi_params, xi_list)
    dU = rho_DmU(theta_DU, xi, k, z, x, pooled) # scalar field at xi
    # partials wrt (kappa, zeta) of the field at xi
    return value_and_grads(dU, (xi.kappa, xi.zeta))
```

Listing 3: Minimal NumPy sketch (DeepSets pooling and readout)

```
import numpy as np

def Phi(params, xi):
    # tiny MLP stub; xi: shape (ds,)
```

```

W1, b1, W2, b2 = params
h = np.tanh(xi @ W1 + b1)          # [h]
return h @ W2 + b2                 # [d]

def embed_measure_np(phi_params, xis):
    # xis: shape [N, ds]; pooled feature: [d]
    feats = np.vstack([Phi(phi_params, xi) for xi in xis]) # [N, d]
    return feats.mean(axis=0)

def rho_U(theta, k, z, x, pooled):
    # simple linear readout for illustration
    W, b = theta
    inp = np.concatenate([np.array([k, z, x]), pooled])
    return float(inp @ W + b)

def U_np(theta_U, phi_params, k, z, x, xis):
    pooled = embed_measure_np(phi_params, xis)
    return rho_U(theta_U, k, z, x, pooled)

# Complexity:  $O(Nd)$  for  $\Phi$ ; pooling is  $O(Nd)$ . Conditioning: scale inputs,
# use Lipschitz activations, and average pooling for stability across  $N$ .

```

Mathematical Insight: Rigor & Implications

Conditioning and invariances (Route B).

- *Permutation invariance:* Average pooling ensures $U(k, z, x, m)$ is invariant to the ordering of atoms in empirical m .
- *Scale/shift:* Standardize (k, z, x) and feature outputs of Φ to improve conditioning; keep readouts Lipschitz.
- *Complexity:* For N atoms and feature width d , evaluating Φ is $\mathcal{O}(Nd)$ and pooling is $\mathcal{O}(Nd)$ per sample.

Listing 4: Pseudo-JAX training loop for Route B (ME residual minimization)

```

import jax, jax.numpy as jnp

def me_residuals(params, batch):
    # batch: list of tuples (k, z, x, xi_list) with empirical measures
    # returns residuals per sample (shape [B]) combining HJB and transport terms
    # NOTE: U, DmU, and transport evaluation are assumed available
    def resid_one(sample):
        k, z, x, xi_list = sample
        # compute U, grads, and transport using DeepSets encodings
        return eval_me_residual(params, k, z, x, xi_list) # scalar residual
    return jax.vmap(resid_one)(batch)

```

```

def loss_fn(params, batch):
    res = me_residuals(params, batch)
    loss_me = jnp.mean(res**2)
    # boundary penalty and gauge anchoring for D_m U
    pen_boundary = boundary_penalty(params, batch)
    pen_anchor = anchor_penalty(params, batch)
    return loss_me + 1e-2*pen_boundary + 1e-3*pen_anchor

@jax.jit
def train_step(params, opt_state, batch):
    loss, grads = jax.value_and_grad(loss_fn)(params, batch)
    updates, opt_state = optimizer.update(grads, opt_state)
    params = optax.apply_updates(params, updates)
    return params, opt_state, loss

# Reproducibility: fix seed, device, and dtype
seed = 0
key = jax.random.PRNGKey(seed)
jax.config.update("jax_enable_x64", True) # prefer float64 for PDE stability
device = jax.devices()[0]
print({"seed": seed, "device": device, "dtype": jnp.float64})

# Batching measures: pad or bucket xi_list to fixed length for vmap/jit
for step, batch in enumerate(data_loader):
    params, opt_state, loss = train_step(params, opt_state, batch)
    if step % 50 == 0:
        print(step, float(loss))

```

Pedagogical Insight: Economic Intuition & Context

Reproducibility hooks.

- Fix RNG seeds and report device (CPU/GPU/TPU) and dtype (float32/64).
- Use padding/bucketing for variable-size empirical measures to keep JIT shapes static.
- In notebooks/CLIs, honor a NOTEBOOK_FAST flag to reduce steps/batch for quick checks.

10 Verification and Diagnostics

Residual norms. For collocation tuples (k, z, x, m) :

$$\mathcal{R}_{\text{HJB}} \equiv r(x) V - \max_i \{\pi + V_k(i - \delta k) + L_z V + L_x V\},$$

$$\mathcal{R}_{\text{FP}} \equiv -\partial_k[(i^* - \delta k), m] + L_z^* m,$$

$$\mathcal{R}_{\text{ME}} \equiv r(x) U - \left(\max_i \{\pi + U_k(i - \delta k) + L_z U + L_x U\} + \int \cdots, m(d\xi) \right).$$

Typical norms: L^2 over collocation points or weighted Sobolev norms. KKT and boundary penalties are added for feasibility; in Route A, measure W_2 drifts between iterations provide a sharp

distributional diagnostic.

Stopping rules. Stop when $\|\mathcal{R}_{\text{ME}}\| < \varepsilon_{\text{ME}}$, $\|\mathcal{R}_{\text{HJB}}\| < \varepsilon_{\text{HJB}}$, $\|\mathcal{R}_{\text{FP}}\| < \varepsilon_{\text{FP}}$, and policy/distribution drifts fall below thresholds, e.g., $\sup |i^{*,(n+1)} - i^{*,(n)}| < 10^{-5}$ and $W_2(m^{(n+1)}, m^{(n)}) < 10^{-4}$.




Figure 5: Placeholder: typical convergence of the master-equation residual.

Sanity checks.

- *No-price-limit case.* If P is flat, the price dependence on m vanishes. Route A and B should collapse to the same frictional-control model without cross effects.
- *Symmetric costs.* Setting $\phi_- = \phi_+$ removes the kink; i^* is linear in $V_k - 1$ everywhere. FP becomes smoother; residuals drop faster.
- *Elasticity sweep.* Under isoelastic demand, η scales the marginal-revenue wedge linearly; recovered investment schedules should contract monotonically in η .

Wasserstein-2 drift diagnostic (practical). For empirical measures with equal weights in 1D (e.g., monitoring the k -marginal), the quadratic Wasserstein distance admits an $\mathcal{O}(N \log N)$ implementation via sorting:

Listing 5: Empirical W_2 in 1D via sorting (equal weights)

```
import numpy as np
```

```
def w2_empirical_1d(xs, ys):
    """Quadratic Wasserstein distance for equal-weight samples.
    xs, ys: arrays of shape [N]. Returns  $W_2$  (not squared).
    """
    xs = np.sort(np.asarray(xs))
    ys = np.sort(np.asarray(ys))
    return np.sqrt(np.mean((xs - ys)**2))
```

```
# Example: monitor drift between successive iterates of the FP solver
W2_drift = w2_empirical_1d(samples_k_t, samples_k_t1)
print(f"W2_drift(k-marginal): {W2_drift:.3e}")
```

For higher dimensions, one can approximate W_2 by projecting onto random 1D directions (sliced Wasserstein) or use optimal-transport solvers (entropic regularization) at higher cost.

Sliced Wasserstein (2D approximation). Project samples onto random unit directions and average 1D distances; complexity $\mathcal{O}(RN \log N)$ for R projections:

Listing 6: Sliced W_2 for 2D (equal weights)

```
import numpy as np

def sliced_w2_2d(X, Y, R=64, rng=None):
    """Approximate  $W_2$  in 2D via random 1D projections.
    X, Y: arrays [N, 2]; returns sliced  $W_2$  (not squared).
    """
    rng = np.random.default_rng(rng)
    X = np.asarray(X); Y = np.asarray(Y)
    assert X.shape == Y.shape and X.shape[1] == 2
    acc = 0.0
    for _ in range(R):
        theta = rng.normal(size=2)
        theta = theta / np.linalg.norm(theta)
        x1d = X @ theta; y1d = Y @ theta
        acc += w2_empirical_1d(x1d, y1d)**2
    return np.sqrt(acc / R)
```

Mathematical Insight: Rigor & Implications

Bias/variance tradeoff. The sliced- W_2 is a lower bound on W_2 ; variance shrinks as R grows. For diagnostics, modest R (e.g., 32–128) often suffices to detect drift trends.

11 Economics: Aggregation, Irreversibility, Comparative Statics

Aggregation. Aggregation enters through $P(Y(m, x))$ within the HJB. Under isoelastic demand, the effective marginal-revenue wedge is $-\eta P(Y) e^{x+z} k^\alpha$, which acts as a proportional reduction in marginal revenue.

Irreversibility. The asymmetry $\phi_- > \phi_+$ creates a kink in the Hamiltonian and investment bands: for V_k just below 1 the disinvestment response is muted relative to the investment response for V_k just above 1. At the distributional level, this slows the left-tail motion in k , thickening the mass near low capital.

Comparative statics.

- Larger η (steeper demand) amplifies the negative externality, reducing investment and shifting mass in m toward lower k .
- Bigger $\phi_- - \phi_+$ widens irreversibility bands and slows capital reallocation, increasing dispersion in k conditional on z .

- Higher σ_z spreads the cross-section in z , raising Y volatility and, through $P'(Y)$, modulating the marginal-revenue wedge over the business cycle.
- Higher σ_x (through L_x) deepens precautionary effects via $r(x)$ and the HJB drift terms, with ambiguous effects on average investment depending on curvature.
- A countercyclical $r(x)$ strengthens the value premium mechanism \tilde{A} la costly reversibility by raising discount rates in recessions precisely when $P'(Y)$ is most negative.

A Appendix A: Derivations and Technical Lemmas

Mathematical Insight: Rigor & Implications

Correction and Addenda. The derivation of the Master Equation is corrected to exclude a separate $\int \delta_m \pi, dm$ term; see ?? and §??. Below we provide compact verification artifacts for the envelope identity and for the functional chain rule used in the population-transport term.

Envelope identity (with depreciation term)

Let $p = V_k$ and $\mathcal{H}(k, \cdot) = \max_i \{\pi + p(i - \delta k)\}$. Then $\partial_p \mathcal{H} = i^*(p) - \delta k$.

Symbolic Check (SymPy)

```
import sympy as sp
i,k,p,phi_p,phi_m,delta = sp.symbols('i k p phi_p phi_m delta', positive=True)
h_p = sp.Rational(1,2)*phi_p*i**2/k
h_m = sp.Rational(1,2)*phi_m*i**2/k
sol_p = k*(p-1)/phi_p
sol_m = k*(p-1)/phi_m
H_p = (-i - h_p + p*(i - delta*k)).subs(i, sol_p)
H_m = (-i - h_m + p*(i - delta*k)).subs(i, sol_m)
assert sp.simplify(sp.diff(H_p,p) - (sol_p - delta*k)) == 0
assert sp.simplify(sp.diff(H_m,p) - (sol_m - delta*k)) == 0
```

Functional chain rule (structure check)

For a one-dimensional diffusion with drift μ and variance σ^2 , the chain rule applied to $D_m F$ has the schematic form $\mu \partial_\xi (D_m F) + \frac{1}{2} \sigma^2 \partial_\xi^2 (D_m F)$.

Symbolic Check (SymPy)

```
import sympy as sp
xi = sp.symbols('xi', real=True)
mu = sp.Function('mu')(xi)
sig2 = sp.symbols('sig2', positive=True)
g = sp.Function('g')(xi)
expr = mu*sp.diff(g, xi) + sp.Rational(1,2)*sig2*sp.diff(g, xi, 2)
assert expr.has(sp.Derivative(g, xi)) and expr.has(sp.Derivative(g, (xi,2)))
```

Lemma A.1: Functional Chain Rule for FP Flows

Let m_t solve the FP equation $\partial_t m = \mathcal{L}^* m$ with drift $b(\xi)$ and diffusion matrix $\sigma(\xi)\sigma(\xi)^\top$. If $F : \mathcal{P}_2(S) \rightarrow \mathbb{R}$ is sufficiently regular (in the sense of Lions), then

$$\frac{d}{dt} F(m_t) = \int_S \left(\langle \nabla_\xi (D_- m F)(m_t)(\xi), b(\xi) \rangle + \frac{1}{2} \text{Tr}[\sigma \sigma^\top(\xi) \nabla_\xi^2 (D_- m F)(m_t)(\xi)] \right) m_t(d\xi).$$

Proof sketch. Apply the functional Itô calculus on \mathcal{P}_2 to $F(m_t)$ and the adjoint pairing to identify the generator acting componentwise on $D_- m F$. See [?, Ch. 5] and [?].

A.1 Envelope/KKT and policy recovery

From (??), define $p = V_k$. The Hamiltonian $\mathcal{H}(k, z, x, m, p) = \max_i \{\pi + p(i - \delta k)\}$ is the convex conjugate of h shifted by $p - 1$. The envelope condition $V_k = \partial_p \mathcal{H}$ combined with the FOC for i produces the piecewise-affine policy in ??. The kink at $p = 1$ corresponds to $i = 0$. KKT adds the complementary slackness $\lambda \cdot (i + \bar{l}(k)) = 0$ when a lower bound is present.

A.2 Adjoint pairing for FP

Let φ be a smooth test function. Then

$$\frac{d}{dt} \int \varphi, dm_t = \int \varphi_k(i^* - \delta k), dm_t + \int L_- z \varphi, dm_t = \int \varphi, d \left(-\partial_k[(i^* - \delta k)m_t] + L_- z^* m_t \right).$$

Stationarity imposes (??) with $\partial_t m = 0$. Reflecting at $k = 0$ eliminates the boundary integral.

A.3 Functional Calculus and the Master Equation

Consider a flow $t \mapsto (K_t, Z_t)$ for the tagged firm following control i_t and a flow of measures $t \mapsto m_t$ solving (??) under the feedback $i^*(\cdot, m_t)$. By functional Itô's lemma for $U(K_t, Z_t, x, m_t)$,

$$\begin{aligned} dU &= U_k dK_t + U_z dZ_t + \frac{1}{2} U_{zz} \sigma_z^2 dt \\ &\quad + (\partial_t U|_m), dt, \\ \partial_t U|_m &= \int \left[(i^*(\xi, x, m) - \delta \kappa) \partial_\kappa (D_- m U)_\kappa + \mu_z(\zeta) \partial_\zeta (D_- m U)_\zeta + \frac{1}{2} \sigma_z^2 \partial_{\zeta\zeta}^2 (D_- m U)_\zeta \right] m(d\xi). \end{aligned}$$

Taking expectations under the pricing measure with short rate $r(x)$ and imposing stationarity yields the stationary master equation (ME) in ??: the own-firm HJB terms and the population-transport term. The dependence of π on m (via $P(Y(m, x))$) is already handled inside the Hamiltonian and does not appear as a separate explicit term.

A.4 Externality term in detail

Write $\pi(k, i, z, x, m) = \Psi(Y(m, x)) \chi(k, z, x) - i - h(i, k) - f$ with $\Psi = P$ and $\chi = e^{x+z} k^\alpha$. Then

$$\delta_- m \pi(m)(\xi) = \Psi'(Y) \chi(k, z, x) \chi(\kappa, \zeta, x),$$

and integration wr.t. m yields $\chi(k, z, x) \Psi'(Y) Y(m, x)$.

B Appendix B: Residual-Loss Template (for implementation)

For a collocation tuple (k, z, x) , an empirical measure $m = \frac{1}{N} \sum_{n=1}^N \delta_{\xi^n}$, and parameterized U_ω, D_-mU_ψ , define

$$\begin{aligned} \hat{Y} &\equiv \frac{1}{N} \sum_{n=1}^N e^{x+\zeta^n} (\kappa^n)^\alpha, \\ \hat{\mathcal{R}}_{\text{ME}} &\equiv r(x) U_\omega \\ &\quad - \max_i \left\{ \pi + (U_\omega)_k (i - \delta k) + L_- z U_\omega + L_- x U_\omega \right\} \\ &\quad - \frac{1}{N} \sum_{n=1}^N \left[(i^*(\xi^n, x, m) - \delta \kappa^n) \partial_\kappa D_-mU_\psi(\xi^n) + \mu_z(\zeta^n) \partial_\zeta D_-mU_\psi(\xi^n) + \frac{1}{2} \sigma_z^2 \partial_{\zeta\zeta}^2 D_-mU_\psi(\xi^n) \right] \end{aligned}$$

Add soft KKT penalties (one-sided around $(U_\omega)_k = 1$), an anchoring penalty $\mathcal{P}_{\text{anchor}} = (\int D_-mU, dm)^2$ to fix the gauge, and boundary regularizers (reflecting $k = 0$, growth at k_{max}). Minimize

$$\mathcal{L} = \mathbb{E} [\hat{\mathcal{R}}_{\text{ME}}^2] + \lambda_{\text{KKT}} \mathcal{P}_{\text{KKT}} + \lambda_{\text{bdry}} \mathcal{P}_{\text{bdry}} + \lambda_{\text{anchor}} \mathcal{P}_{\text{anchor}}.$$

Anchoring removes the gauge freedom in D_-mU .

C Appendix C: Common-Noise Master Equation (Reference Note)

When the population law m_t itself diffuses under common noise (say through an exogenous x_t or an aggregate Brownian component shared by firms), the functional Itô calculus on \mathcal{P}_2 introduces a second-order term in the measure variable. In a stylized form (see Carmona & Delarue, and Cardaliaguet–Delarue–Lasry–Lions), the stationary master equation would add a term of the form

$$\frac{1}{2} \Sigma_{\text{com}} : \int \int \partial_\xi \partial_{\xi'} (D_-mU(\xi)) (D_-mU(\xi')) m(\cdot, d\xi) m(\cdot, d\xi')$$

or, in classical PDE notation, $\frac{1}{2} \text{Tr} [\Gamma \partial_{\xi\xi}^2 D_-mU]$ integrated against m , where Γ is the covariance of the common noise. Because this paper conditions on x , these terms are absent in the stationary master equation.

Mathematical Insight: Rigor & Implications

Displacement monotonicity (context). For master equations with common noise, uniqueness and well-posedness often require *displacement monotonicity*: a convexity-type condition along Wasserstein geodesics for the coupling. See [?] for precise statements. In our economic environment, couplings of the type $m \mapsto P(\int q, dm)$ satisfy Lasry–Lions monotonicity when $P'(\cdot) < 0$ (??), but displacement monotonicity may require additional curvature restrictions on P or on the state-cost structure. Since we work conditional on x (no common-noise measure diffusion), we do not invoke displacement monotonicity in this paper.

D Appendix D: Tiny Pseudocode (Plain listings)


```

# Inputs:

# params\_omega: parameters for  $U(k, z, x; m)$ 

# params\_psi: parameters for  $\delta_m U(x_i; k, z, x; m)$ 

# batch: list of tuples  $(k, z, x, \{x_i\}_{n=1}^N = \{\kappa_n, \zeta_n\})$ 

# primitives:  $\alpha, \delta, \mu_z(z), \sigma_z, \mu_x(x), \sigma_x, r(x)$ ,

# demand  $P(Y)$ , fixed cost  $f$ 

# penalties: lambdas for KKT, boundary, and anchor (gauge) regularizers

def policy\_from\_grad(p, k, phi\_plus, phi\_minus):
    \#  $p = U_k$  (value gradient)
    if p >= 1.0:
        return (k/phi\_plus)*(p - 1.0)
    else:
        return (k/phi\_minus)*(p - 1.0)

def reflecting\_penalty(k, i\_star):
    \# discourage negative control at  $k=0$  and large negative flux
    pen0 = max(0.0, -i\_star) if k <= 1e-10 else 0.0
    return pen0\**2

def h\_cost(i, k, phi\_plus, phi\_minus):
    if i >= 0.0:
        return 0.5*phi\_plus*(i*i)/max(k, 1e-12)
    else:
        return 0.5*phi\_minus*(i*i)/max(k, 1e-12)

def HJB\_operator(k, z, x, Yhat, Uk, Uz, Uzz, Ux, Uxx, i):
    q = exp(x+z)*(k\**alpha)
    pi = P(Yhat)*q - i - h\_cost(i, k, phi\_plus, phi\_minus) - f
    return pi + Uk*(i - delta*k) + mu\_z(z)*Uz + 0.5*sigma\_z\**2*Uzz\**20;
    \+ mu\_x(x)*Ux + 0.5*sigma\_x\**2*Uxx

def ME\_residual\_for\_tuple(params\_omega, params\_psi, tup):
    k, z, x, xi\_list = tup.k, tup.z, tup.x, tup.xi\_list
    \# empirical measure moments
    Y\_hat = mean([exp(x+xi.zeta)*(xi.kappa\**alpha) for xi in xi\_list])
    \#  $U$  and its partials at  $(k, z, x)$ 
    U, Uk, Uz, Uzz, Ux, Uxx = U\_and\_grads(params\_omega, k, z, x, xi\_list)
    \# best response  $i^*$ 
    i\_star = policy\_from\_grad(Uk, k, phi\_plus, phi\_minus)
    \# HJB maximand at  $i^*$ 
    H\_val = HJB\_operator(k, z, x, Y\_hat, Uk, Uz, Uzz, Ux, Uxx, i\_star)
    \# Population transport term (via the Lions derivative)
    integ = 0.0
    for xi in xi\_list:
        dU = delta\_mU\_and\_partials(params\_psi, xi, k, z, x, xi\_list)

```

```

\# dU returns dict with fields dkappa, dzeta, dzeta2, p\_k (proxy gradient)
i\_star\_xi = policy\_from\_grad(dU['p\_k'], xi.kappa, phi\_plus, phi\_minus)
integ += (i\_star\_xi - delta*xi.kappa)* dU['dkappa']&#x20;
\+ mu\_z(xi.zeta)\* dU['dzeta']&#x20;
\+ 0.5*sigma\_z\*\*2 \* dU['dzeta2']
integ = integ / len(xi\_list)
\# assemble residual (no separate externality term; price enters via P(Yhat) in HJB)
res = r(x)\*U - max(H\_val, HJB\_operator(k,z,x,Y\_hat,Uk,Uz,Uzz,Ux,Uxx,0.0))&#x20;\-
    integ
\# penalties
pen = reflecting\_penalty(k, i\_star)
return res, pen

def loss(params\_omega, params\_psi, batch):
sse = 0.0
pen = 0.0
for tup in batch:
res, p = ME\_residual\_for\_tuple(params\_omega, params\_psi, tup)
sse += res\*\*2
pen += p
anchor = anchor\_penalty(params\_psi, batch) # e.g., squared mean of dmU over batch
return sse/len(batch) + lambda\_bdry\*pen + lambda\_anchor\*anchor

```

Listing 7: Pseudo-JAX for (ME) residual with empirical measure

E Appendix E: Symbolic Verification (PythonTeX + SymPy)

This appendix runs minimal SymPy checks to verify key derivations used in the text. Compilation is configured (via `latexmkrc`) to execute these checks on every build; any failure triggers a build error. We assume smoothness and reflecting/no-flux boundary conditions where noted.

Symbolic Check (SymPy)

```

import sympy as sp

# 1) Isoelastic simplification:  $Y P'(Y) = -\eta P(Y)$ 
Y, eta = sp.symbols('Y eta', positive=True)
P = Y**(-eta)
check1 = sp.simplify(Y*sp.diff(P, Y) + eta*P)
assert check1 == 0
print("Isoelastic:  $Y P'(Y) = -\eta P(Y)$  [OK]")

# 2) Externality directional derivative:  $\frac{d}{d \epsilon} P(Y + \epsilon \chi_0) \chi_0 \big|_{\epsilon=0}$ 
# equals  $P'(Y) * \chi_0$ 
chi0, chieps, eps = sp.symbols('chi0 chieps eps', real=True)
Psi = lambda y: y**(-eta)
dpi = sp.diff(Psi(Y + eps*chieps)*chi0, eps).subs(eps, 0)
target = sp.diff(Psi(Y), Y) * chi0
assert sp.simplify(dpi - target) == 0
print('Externality directional derivative [OK]')

# 3) Externality, isoelastic reduction after integrating over m:  $\chi_0 * Y * P'(Y) = -\eta * P(Y) * \chi_0$ 
lhs = chi0 * Y * sp.diff(Psi(Y), Y)

```

```

rhs = -eta * Psi(Y) * chi0
assert sp.simplify(lhs - rhs) == 0
print('Externality isoelastic reduction    [OK]')

# 4) KKT/FOC solution for i* with asymmetric quadratic costs
#   h = 0.5*phi_plus*i^2/k for i>=0;   0.5*phi_minus*i^2/k for i<0
i, k, p, phi_plus, phi_minus = sp.symbols('i k p phi_plus phi_minus', positive=True)
h_plus = 0.5*phi_plus*i**2/k
FOC_plus = sp.Eq(sp.diff(-i - h_plus + p*i, i), 0)
sol_plus = sp.solve(FOC_plus, i)[0]
h_minus = 0.5*phi_minus*i**2/k
FOC_minus = sp.Eq(sp.diff(-i - h_minus + p*i, i), 0)
sol_minus = sp.solve(FOC_minus, i)[0]
assert sp.simplify(sol_plus - k*(p-1)/phi_plus) == 0
assert sp.simplify(sol_minus - k*(p-1)/phi_minus) == 0
print('KKT/FOC piecewise i* formulas    [OK]')

# 5) FP adjoint pairing identity (algebraic, boundary terms omitted):
#   phi_k * (a*m) = d_k(phi*a*m) - phi * d_k(a*m)
kk = sp.symbols('kk', real=True)
phi = sp.Function('phi')(kk)
a = sp.Function('a')(kk)
mm = sp.Function('m')(kk)
expr = sp.diff(phi, kk)*(a*mm) - (sp.diff(phi*(a*mm), kk) - phi*sp.diff(a*mm, kk))
assert sp.simplify(expr) == 0
print('Adjoint pairing identity (no-flux) [OK]')

# 6) Envelope property for Hamiltonian in p: d/dp max_i { -i - h(i,k) + p i } = i*(p)
#   Check separately on each branch (ignoring terms not depending on i, e.g., -delta*k*p)
H_plus = (-i - h_plus + p*i).subs(i, sol_plus)
H_minus = (-i - h_minus + p*i).subs(i, sol_minus)
dHp_dp = sp.simplify(sp.diff(H_plus, p))
dHm_dp = sp.simplify(sp.diff(H_minus, p))
assert sp.simplify(dHp_dp - sol_plus) == 0
assert sp.simplify(dHm_dp - sol_minus) == 0
print('Envelope: dH/dp equals i*(p)    [OK]')

print('\nAll SymPy verification checks passed.')

```

F Appendix F: Lean4 Micro-Proofs (Sketches)

The following Lean4/mathlib4 snippets formalize two identities used in the text. They are provided as self-contained, runnable sketches (assuming a recent mathlib4): the isoelastic simplification $Y P'(Y) = -\eta P(Y)$ and the algebraic reduction $Y \cdot Y^{-\eta-1} = Y^{-\eta}$ for $Y > 0$.

Formal Proof (Lean4)

```

import Mathlib.Analysis.Calculus.Deriv
import Mathlib.Data.Real.Basic

open Real

```

```

variable {Î· Y : R}

-- P(Y) = Y ^ (-Î·), defined for Y > 0 via rpow
def P (Y : R) (Î· : R) : R := Y ^ (-Î·)

-- Algebraic reduction: for Y > 0, Y * Y^(-Î· - 1) = Y^(-Î·)
theorem rpow_mul_cancel (hY : 0 < Y) :
  Y * Y ^ (-Î· - 1) = Y ^ (-Î·) := by
  -- rewrite Y as Y^1 and use rpow_add (valid for Y > 0)
  have h1 : Y = Y ^ (1 : R) := by simp using (rpow_one Y)
  calc
    Y * Y ^ (-Î· - 1)
      = Y ^ (1 : R) * Y ^ (-Î· - 1) := by simp [h1]
    _ = Y ^ ((1 : R) + (-Î· - 1)) := by
      simp using (rpow_mul_rpow_of_pos hY (1 : R) (-Î· - 1))
    _ = Y ^ (-Î·) := by ring

-- Differential identity: for Y > 0, (Y) * (deriv (fun y => P y Î·) Y) = -Î· * P Y Î·
theorem isoelastic_identity (hY : 0 < Y) :
  Y * (deriv (fun y => P y Î·) Y) = -Î· * P Y Î· := by
  -- mathlib: d/dy (y^a) = a * y^(a-1) for y>0
  have hderiv : deriv (fun y => y ^ (-Î·)) Y = (-Î·) * Y ^ (-Î· - 1) := by
    simp using (deriv_rpow_const (x:=Y) (a:=-Î·) hY.ne')
  -- multiply both sides by Y and reduce
  calc
    Y * (deriv (fun y => P y Î·) Y)
      = Y * ((-Î·) * Y ^ (-Î· - 1)) := by simp [P, hderiv]
    _ = -Î· * (Y * Y ^ (-Î· - 1)) := by ring
    _ = -Î· * Y ^ (-Î·) := by simp using rpow_mul_cancel (Î·:=Î·) (Y:=Y) hY
    _ = -Î· * P Y Î· := by rfl

```

Pedagogical Insight: Economic Intuition & Context

Notes. The lemmas use `rpow` and standard calculus from `mathlib4`. They require $Y > 0$ for real-exponent laws. The SymPy checks in `??` independently validate the same identities numerically/symbolically.

Formal Proof (Lean4)

```

import Mathlib.Data.Real.Basic

-- Risk coefficient appearing in the EZ market price of risk
def risk_coeff (Î³ Î· : R) : R := (1 - Î³) * (1 - 1/Î·)

-- If RRA = 1 (log utility), the utility-channel risk coefficient vanishes
@[simp] lemma risk_coeff_gamma_one (Î· : R) : risk_coeff 1 Î· = 0 := by
  simp [risk_coeff]

-- If EIS = 1, the utility-channel risk coefficient vanishes
@[simp] lemma risk_coeff_psi_one (Î³ : R) : risk_coeff Î³ 1 = 0 := by
  simp [risk_coeff]

```

G Appendix G: Endogenous SDF with Epstein–Zin Aggregator

When the stochastic discount factor (SDF) is endogenous, it is often derived from a representative agent’s preferences. Epstein–Zin (EZ) preferences allow separating the elasticity of intertemporal substitution (EIS) from relative risk aversion (RRA), which is crucial for asset pricing implications. This appendix details the continuous-time EZ aggregator used as a BSDE driver and derives the corresponding pricing kernel exposure.

Definition G.1: Continuous-time Epstein–Zin Aggregator

Fix time preference $\varrho > 0$ (using notation from ??), risk aversion $\gamma > 0$, and elasticity of intertemporal substitution $\psi > 0$. We assume $\psi \neq 1$. Let

$$\vartheta \equiv \frac{1 - \gamma}{1 - 1/\psi}.$$

The parameter ϑ is sometimes used in alternative normalizations of the utility index. For aggregate consumption $c_t > 0$, the representative agent’s utility process (V_t, Z_t) solves the backward SDE:

$$, dV_t = -f(c_t, V_t, Z_t), dt + Z_t^\top, dB_t,$$

where $V_t > 0$ is the continuation value and $Z_t \in \mathbb{R}^{d_B}$ is the exposure vector to aggregate Brownian shocks B_t . The driver f is the EZ aggregator. We use the standard normalization (consistent with Duffie–Epstein, 1992, [?], and also used in [?]):

$$f(c, V, Z) = \frac{\varrho}{1 - 1/\psi} \left(c^{1-1/\psi} V^{1/\psi} - V \right) + \frac{1}{2} (1 - \gamma) (1 - 1/\psi) \frac{\|Z\|^2}{V}. \quad (\text{G.1})$$

Pedagogical Insight: Economic Intuition & Context

Economic Intuition: Separating RRA and EIS. EZ preferences break the link imposed by standard CRRA utility (where $\text{EIS} = 1/\text{RRA}$).

- γ (RRA) controls aversion to static risk (gambles).
- ψ (EIS) controls willingness to substitute consumption over time (smoothness).

The aggregator f includes an intertemporal substitution term (first part) and a risk adjustment term (second part). The sign of the risk adjustment coefficient $(1 - \gamma)(1 - 1/\psi)$ determines the preference for early ($\gamma > 1, \psi > 1$ or $\gamma < 1, \psi < 1$) or late resolution of uncertainty.

Proposition G.1: Pricing kernel exposure under EZ

Let M_t denote the stochastic discount factor. The utility-channel diffusion component of the instantaneous market price of risk implied by ?? is

$$\Lambda_t^{\text{util}} = \partial_Z f(c_t, V_t, Z_t) = (1 - \gamma) (1 - 1/\psi) \frac{Z_t}{V_t},$$

entering $dM_t/M_t = -r_t dt - (\Lambda_t)^\top dB_t$. If consumption c_t carries its own Brownian exposure, the total Λ_t adds the consumption channel in the usual way.

Proof. The pricing kernel M_t ensures that the utility process V_t , when discounted by M_t , is a martingale: $\mathbb{E}_t[M_{t+s}V_{t+s}] = M_tV_t$. Applying Itô's lemma to the product M_tV_t gives

$$,d(M_tV_t) = V_t, dM_t + M_t, dV_t + \langle, dM_t, , dV_t \rangle.$$

Substitute the dynamics

$$\begin{aligned} ,dM_t/M_t &= -r_t, dt - \Lambda_t^\top, dB_t, \\ ,dV_t &= -f(c_t, V_t, Z_t), dt + Z_t^\top, dB_t, \end{aligned}$$

and note the quadratic covariation $\langle, dM_t, , dV_t \rangle = -M_t \Lambda_t^\top Z_t, dt$. For M_tV_t to be a martingale, the drift must vanish:

$$V_t(-M_t r_t) + M_t(-f(c_t, V_t, Z_t)) - M_t \Lambda_t^\top Z_t = 0,$$

yielding the generalized HJB identity $r_t V_t + f(c_t, V_t, Z_t) + \Lambda_t^\top Z_t = 0$. In equilibrium, the utility-channel component of the market price of risk is identified with the gradient of the aggregator with respect to Z_t (see [?]). Differentiating ?? w.r.t. Z gives

$$\partial_Z f(c, V, Z) = \frac{1}{2}(1 - \gamma)(1 - 1/\psi) \partial_Z \left(\frac{\|Z\|^2}{V} \right) = (1 - \gamma)(1 - 1/\psi) \frac{Z}{V},$$

which proves the claim. □

Symbolic Check (SymPy)

```
import sympy as sp
# Verify the gradient calculation in the proof of Proposition G.1
# using an element-wise approach for robustness.
c, V, gamma, psi = sp.symbols('c V gamma psi', positive=True)
z1, z2 = sp.symbols('z1 z2', real=True)
Z = sp.Matrix([z1, z2])

# Coefficient in the standard normalization (Duffie--Epstein)
coeff_risk = (1-gamma)*(1-1/psi)

# Risk term: (1/2) * coeff * (Z^T Z) / V
risk_term = sp.Rational(1, 2) * coeff_risk * (Z.T * Z)[0, 0] / V

# Gradient with respect to Z's components
grad_Z = sp.Matrix(sp.derive_by_array(risk_term, [z1, z2]))
expected = coeff_risk * Z / V

# Robust check: ensure both components are exactly zero
difference = sp.simplify(grad_Z - expected)
assert all(e == 0 for e in difference)
```

Mathematical Insight: Rigor & Implications

Integration with the Firm Problem. When using an endogenous SDF, the firm's HJB equation (??) incorporates the market price of risk Λ_t :

$$r_t V = \max_i \left\{ \pi + V_k(i - \delta k) + L_- z V + L_- x V - (\sigma_z V_z, \sigma_x V_x) \cdot \Lambda_t \right\}.$$

If the aggregate shocks x correspond to the Brownian motion B_t driving EZ utility, this links firm valuation to representative-agent preferences via Λ_t .

Pedagogical Insight: Economic Intuition & Context

Implementation hook. The repository exposes a JAX-friendly generator for ?? and a utility-channel SDF exposure helper:

bsde_dsgE/models/epstein_zin.py: EZParams, ez_generator, sdf_exposure_from_ez
bsde_dsgE/models/multicountry.py: preference="EZ" to enable the aggregator

Usage sketch in code:

```
from bsde_dsgE.models.epstein_zin import EZParams
from bsde_dsgE.models.multicountry import multicountry_model

params = EZParams(rho=0.02, gamma=10.0, psi=1.5) # Use rho for time preference
problem = multicountry_model(dim=5, preference="EZ", ez_params=params)
```

The consumption mapping $c_{\text{fn}}(x)$ can be provided by the user; by default, the model uses a positive aggregator from dividend-like states.

References

- [1] Carmona, R. and F. Delarue (2018). *Probabilistic Theory of Mean Field Games with Applications*. Springer.
- [2] Cardaliaguet, P., F. Delarue, J.-M. Lasry, and P.-L. Lions (2019). *The Master Equation and the Convergence Problem in Mean Field Games*. Princeton University Press.
- [3] Lasry, J.-M. and P.-L. Lions (2007). Mean field games. *Japanese Journal of Mathematics* 2(1): 229–260.
- [4] Zaheer, M., S. Kottur, S. Ravanbakhsh, B. Poczos, R. Salakhutdinov, and A. Smola (2017). Deep Sets. *Advances in Neural Information Processing Systems (NeurIPS)* 30.
- [5] Mou, C.-H. and J. Zhang (various years). Second-order master equations with common noise and displacement monotonicity. (Working papers / journal articles; see also related notes by Gangbo–Mou–Zhang.)
- [6] Zhang, L. (2005). The value premium. *Journal of Finance* 60(1): 67–103.
- [7] Duffie, D. and L. G. Epstein (1992). Stochastic Differential Utility. *Econometrica* 60(2): 353–394.
- [8] Sauzet, M. (2023). Recursive preferences in continuous time and implications for asset pricing. (Working paper; aggregator normalisation consistent with ??.)

Symbol	Type	Meaning
<i>States, Controls, and Shocks</i>		
k	state	Capital (≥ 0); reflecting boundary at $k = 0$
i	control	Net investment; $dk = (i - \delta k), dt$
z	state	Idiosyncratic productivity; diffusion with generator L_z
x	state	Aggregate (business-cycle) shock; generator L_x
σ_z, σ_x	parameters	Diffusion volatilities of z and x
μ_z, μ_x	functions	Drift coefficients in L_z, L_x
W, B	processes	Brownian motions for z and x (independent)
<i>Technology and Market Primitives</i>		
$q(k, z, x)$	output	$e^{x+z} k^\alpha, \alpha \in (0, 1)$
$P(\cdot)$	function	Inverse demand; $P' = P'(Y) < 0$
α	parameter	Capital elasticity in production
δ	parameter	Depreciation rate
$h(i, k)$	function	Irreversible adjustment cost (convex, asymmetric)
ϕ_\pm	parameters	Adjustment-cost curvatures for $i \gtrless 0$
f	parameter	Fixed operating cost
η	parameter	Demand elasticity for isoelastic $P(Y) = Y^{-\eta}$
$r(x)$	function	Short rate (or constant ρ) under pricing measure
<i>Measure Theory and Operators</i>		
S	space	State space $\mathbb{R}_+ \times \mathbb{R}$ for (k, z)
m	measure	Cross-sectional law on S
$\mathcal{P}_2(S)$	space	Probability measures on S with finite second moments
W_2	metric	Quadratic Wasserstein distance on $\mathcal{P}_2(S)$
$\xi = (\kappa, \zeta)$	point	Generic element in support of m ("marginal firm")
D_m	operator	Lions derivative operator (measure Fréchet derivative)
$D_m U(\xi; k, z, x, m)$	function	Lions derivative of $m \mapsto U(k, z, x, m)$ at ξ
L_z, L_x	operators	Generators in z and x ; L_z^* is the adjoint of L_z
\mathcal{T}	operator	Transport operator acting on $D_m U$ in (ME)
<i>Equilibrium Objects</i>		
$\pi(\cdot)$	function	Dividends $P(Y)e^{x+z}k^\alpha - i - h(i, k) - f$
$Y(m, x)$	scalar	Aggregate quantity $\int e^{x+z}k^\alpha m(\cdot, dk, \cdot, dz)$
$V(k, z, x; m)$	function	Stationary value function (HJB)
$U(k, z, x, m)$	function	Master value function (ME)
$i^*(\cdot)$	policy	Optimal net investment from HJB/KKT
$\bar{u}(k)$	function	Lower bound on disinvestment (optional)
<i>Representative-agent block (endogenous SDF)</i>		
γ	parameter	Relative risk aversion (RRA) in Epstein–Zin preferences
ψ	parameter	Elasticity of intertemporal substitution (EIS)
ϑ	parameter	Preference index $\vartheta = (1 - \gamma)/(1 - 1/\psi)$
ϱ	parameter	Subjective discount rate (avoids clash with depreciation δ)
M_t	process	Stochastic discount factor (pricing kernel)
r_t	process	Real short rate implied by M_t
Λ_t	process	Market price of risk (Brownian exposure of M_t)

Table 1: Notation used throughout.

Component	Weight	Notes
ME residual MSE	1.0	$\mathcal{L}_{\text{ME}} = \mathbb{E}[\mathcal{R}_{\text{ME}}^2]$; primary objective.
Boundary penalty	10^{-2}	Enforce reflecting boundary and admissibility constraints.
KKT penalty	10^{-2}	Complementarity for investment constraints (if active).
Gauge anchor	10^{-3}	Fix $\int D_- m U, dm$ to remove invariance.

Table 2: Route B loss components and typical starting weights. Tune per calibration and scale of residuals.

Residual	Tight	Medium	Coarse
ε_{ME}	10^{-5}	10^{-4}	10^{-3}
ε_{HJB}	10^{-7}	10^{-6}	10^{-5}
ε_{FP}	10^{-7}	10^{-6}	10^{-5}

Table 3: Suggested tolerances (dimensionless; scale to data).